Propensity Score Analysis in R - Basics

Statistical Horizons

September 15-17, 2022

Contents

1 Day 1 - September 15, 2022 **3**
1.1 Greedy Nearest Neighbor Matching . . . 3
1.1.1 Load Packages . . . 3
1.1.2 Description of Dataset . . . 3
1.1.3 Load Data and Sort . . . 3
1.1.4 Check Balance Before Matching . . . 3
1.1.5 Greedy Nearest Neighbor Matching Without Replacement . . . 7
1.1.5.1 Check Common Support . . . 8
1.1.5.2 Check Balance . . . 9
1.1.6 Greedy Nearest Neighbor Mahalanobis Distance Matching Without Replacement . . . 12
1.1.6.1 Check Balance . . . 12
1.2 Propensity Score Estimation Using Generalized Boosted Regression . . . 15
1.2.1 Load Package . . . 15
1.2.2 Load Data and Sort . . . 15
1.2.3 Fit Generalized Boosted Regression Model . . . 15
1.2.4 Estimate Propensity Scores . . . 16
1.2.5 Plot Propensity Score Distributions . . . 17
1.2.6 Summary Statistics of Propensity Scores . . . 17
1.2.7 GBR Using the WeightIt Package . . . 18
1.2.7.1 Check Balance . . . 19

2 Day 2 - September 16, 2022 **20**
2.1 Optimal Matching . . . 20
2.1.1 Load Packages . . . 20
2.1.2 Load Data . . . 20
2.1.3 Bivariate Comparisons . . . 20
2.1.4 Outcome Models Without Bias Control . . . 23
2.1.5 Matching . . . 26
2.1.5.1 Distribution of estimated propensity scores . . . 26
2.1.5.2 Optimal Matching . . . 28
2.1.5.3 Covariate Imbalance Before and After Matching by Matching Scheme . . . 31
2.1.6 Post-Full-Matching Analysis of Outcome . . . 32
2.1.6.1 Hodges-Lehmann Aligned Rank Test . . . 32
2.1.6.2 Cohen's d . . . 32
2.1.7 Post-Matching Analysis Using Regression of Difference Scores . . . 32
2.1.8 Optimal Matching Using the MatchIt Package . . . 33
2.2 Propensity Score Weighting . . . 34
2.2.1 Load Packages . . . 34
2.2.2 Description of Dataset . . . 34
2.2.3 Estimate ATE and ATT Weights . . . 34

2.2.4 Load Data with Propensity Scores and Calculate Weights . . . 35
2.2.5 Calculate Weights with the WeightIt Package . . . 35
2.2.6 Outcome Analysis . . . 35
2.2.6.1 Weighted Regression with ATE Weights . . . 35
2.2.6.2 Weighted Regression with ATT Weights . . . 36
2.2.7 Check Balance . . . 36

3 Day 3 - September 17, 2022 **39**
3.1 Matching Estimators . . . 39
3.1.1 Load Packages . . . 39
3.1.2 Description of Dataset . . . 39
3.1.3 Load Data . . . 40
3.1.4 Breusch-Pagan Test for Heteroskedasticity . . . 40
3.1.5 Matching Estimators . . . 41
3.1.6 Define Outcome (Y), Treatment Index (Tr), and Variables to Match On (X) . . . 42
3.1.7 Get Estimators Individually . . . 42
3.1.8 Get All Estimators . . . 43
3.1.9 Specify Variables in the Bias-Corrected Matching . . . 44
3.1.10 Check Balance . . . 45

4 Appendix A: IMBALANCE Stata Module in R **47**
4.1 Example Usage . . . 47

5 Appendix B: HODGESL Stata Module in R **48**

1 Day 1 - September 15, 2022

1.1 Greedy Nearest Neighbor Matching

1.1.1 Load Packages

The `haven` and `sjlabelled` packages are used to load and clean Stata data files (.dta); the `MatchIt` package contains functions for greedy matching[1]; the `cobalt` package contains functions for balance checking; and the `tidyverse` package is loaded for its data manipulation functions.

```
library(haven)
library(sjlabelled)
library(cobalt)
library(MatchIt)
library(tidyverse)
library(kableExtra)
library(survival)
```

1.1.2 Description of Dataset

This dataset is a sample of 2,758 children from the National Survey of Child and Adolescent Well-Being (NSCAW). The treatment condition is `aodserv` or caregivers who received (aodserv = 1) or did not receive (aodserv = 0) substance abuse services. Two matching procedures are illustrated here. In Section 5.8.1 of the PSA-R code, 12 matching schemes are shown.

1.1.3 Load Data and Sort

When non-treated cases have the same propensity score values, their matches will depend on the order of the data. Therefore, it's important to order the observations randomly.

```
set.seed(1000)
gm_df <- haven::read_dta("data/chpt5_1_original.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble() %>%
  add_column(runif = runif(nrow(.))) %>%
  arrange(runif) %>%
  select(-runif)
```

1.1.4 Check Balance Before Matching

Because all of the variables that will be used for predicting the propensity scores of service receipt are categorical, we can use `chisq.test` to check their balance before matching.

```
# Change categorical independent variables to factor variables
gm_df <- gm_df %>%
  mutate(across(c(-id, -aodserv), as.factor))

# Marital Status (large sample size, therefore continuity correction not needed)
chisq.test(gm_df$aodserv, gm_df$married, correct = F)
```

```
##
##  Pearson's Chi-squared test
##
## data:  gm_df$aodserv and gm_df$married
## X-squared = 2.9705, df = 1, p-value = 0.0848
```

[1] https://kosukeimai.github.io/MatchIt/articles/matching-methods.html

```r
# All Variables
gm_df %>%
  select(
    married, educ, pov, employ, open, race, chdage, cgage, CRA47A, mental,
    arrest, PSH17A, maltx, ra, cidi, cgneed, cwwrep, aodserv
  ) %>%
  pivot_longer(-aodserv, names_to = "variable") %>%
  group_by(variable) %>%
  nest() %>%
  mutate(bivariate.test = map(data, ~ chisq.test(.$aodserv, .$value, correct = F))) %>%
  mutate(statistic = map(bivariate.test, ~ round(.$statistic, 3))) %>%
  mutate(p.value = map(bivariate.test, ~ round(.$p.value, 3))) %>%
  unnest(cols = c(statistic, p.value)) %>%
  select(variable, statistic, p.value)
```

```
## # A tibble: 17 x 3
## # Groups:   variable [17]
##    variable statistic p.value
##    <chr>        <dbl>   <dbl>
##  1 married       2.97   0.085
##  2 educ         10.5    0.005
##  3 pov          11.3    0.023
##  4 employ       23.1    0
##  5 open         58.4    0
##  6 race         11.5    0.009
##  7 chdage       55.4    0
##  8 cgage         3.56   0.313
##  9 CRA47A       17.5    0
## 10 mental       92.5    0
## 11 arrest      127.     0
## 12 PSH17A      179.     0
## 13 maltx        49.7    0
## 14 ra          585.     0
## 15 cidi        157.     0
## 16 cgneed      139.     0
## 17 cwwrep     1240.     0
```

Alternatively, the `cobalt` package provides several convenient functions for assessing balance.

The standardized mean difference (SMD) is a commonly used balance measure.[2] It is calculated as the difference in means of a covariate across the treatment groups, divided by the standard deviation in the treated group (ATT), the control group (ATC), or the pooled standard deviation (ATE). See `?cobalt::col_w_smd` for additional options. Stuart et al. (2013) recommend 0.1 or 0.25 as reasonable cut-offs for acceptable standardized biases.[3]

```r
# Balance table
cobalt::bal.tab(select(
  gm_df, married, educ, pov, employ, open, race, chdage, cgage, CRA47A,
  mental, arrest, PSH17A, maltx, ra, cidi, cgneed, cwwrep
),
```

[2]The disadvantage of hypothesis tests is "they are influenced by sample size, which fluctuates during adjustment, and the theory behind them is inappropriate because balance is a quality solely of the sample in question, not in relation to a population" (https://cran.r-project.org/web/packages/cobalt/vignettes/cobalt.html).

[3]Stuart, E. A., Lee, B. K., & Leacy, F. P. (2013). Prognostic score–based balance measures for propensity score methods in comparative effectiveness research. *Journal of Clinical Epidemiology*, 66(8 0), S84-S90.e1. https://doi.org/10.1016/j.jclinepi.2013.01.013

```r
treat = gm_df$aodserv,
s.d.denom = "treated",
threshold = .1,
stats = "mean.diff",
continuous = "std",
binary = "std"
)
```

```
## Balance Measures
##            Type Diff.Un     M.Threshold.Un
## married  Binary -0.1108 Not Balanced, >0.1
## educ_0   Binary  0.1605 Not Balanced, >0.1
## educ_1   Binary -0.0236     Balanced, <0.1
## educ_2   Binary -0.1840 Not Balanced, >0.1
## pov_1    Binary  0.1395 Not Balanced, >0.1
## pov_2    Binary  0.0628     Balanced, <0.1
## pov_3    Binary -0.1584 Not Balanced, >0.1
## pov_4    Binary -0.0829     Balanced, <0.1
## pov_5    Binary -0.0410     Balanced, <0.1
## employ   Binary -0.3082 Not Balanced, >0.1
## open     Binary  0.5516 Not Balanced, >0.1
## race_1   Binary -0.1167 Not Balanced, >0.1
## race_2   Binary  0.0807     Balanced, <0.1
## race_3   Binary -0.0519     Balanced, <0.1
## race_4   Binary  0.1392 Not Balanced, >0.1
## chdage_0 Binary -0.2527 Not Balanced, >0.1
## chdage_1 Binary  0.4279 Not Balanced, >0.1
## chdage_2 Binary -0.0856     Balanced, <0.1
## chdage_3 Binary -0.2473 Not Balanced, >0.1
## cgage_0  Binary  0.0781     Balanced, <0.1
## cgage_1  Binary -0.0694     Balanced, <0.1
## cgage_2  Binary  0.0300     Balanced, <0.1
## cgage_3  Binary  0.0297     Balanced, <0.1
## CRA47A   Binary  0.2406 Not Balanced, >0.1
## mental   Binary  0.5238 Not Balanced, >0.1
## arrest   Binary  0.6718 Not Balanced, >0.1
## PSH17A   Binary  0.5347 Not Balanced, >0.1
## maltx_0  Binary -0.2025 Not Balanced, >0.1
## maltx_1  Binary -0.4690 Not Balanced, >0.1
## maltx_2  Binary  0.2909 Not Balanced, >0.1
## maltx_3  Binary  0.0690     Balanced, <0.1
## maltx_4  Binary  0.0110     Balanced, <0.1
## ra       Binary  1.1915 Not Balanced, >0.1
## cidi     Binary  0.7110 Not Balanced, >0.1
## cgneed   Binary  0.3907 Not Balanced, >0.1
## cwwrep   Binary  1.6213 Not Balanced, >0.1
##
## Balance tally for mean differences
##                     count
## Balanced, <0.1         13
## Not Balanced, >0.1     23
##
## Variable with the greatest mean difference
##  Variable Diff.Un     M.Threshold.Un
```

```
##    cwwrep  1.6213 Not Balanced, >0.1
##
## Sample sizes
##     Control Treated
## All    2460     298
```

```
# Love plot
cobalt::love.plot(select(
  gm_df, married, educ, pov, employ, open, race, chdage, cgage, CRA47A,
  mental, arrest, PSH17A, maltx, ra, cidi, cgneed, cwwrep
),
treat = gm_df$aodserv,
binary = "std",
s.d.denom = "treated",
threshold = .1,
) +
  labs(title = "Covariate Balance Before Matching")
```

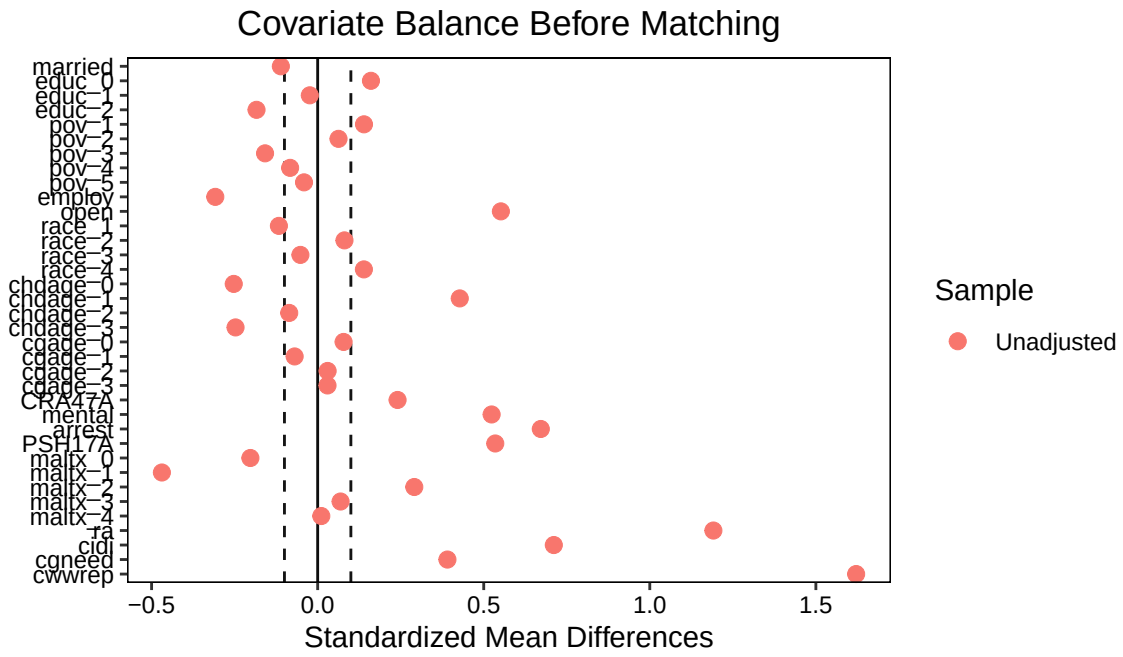

1.1.5 Greedy Nearest Neighbor Matching Without Replacement

By default, the `MatchIt::matchit()` function performs greedy nearest neighbor matching without replacement, therefore the `method = "nearest"` and `replace = F` arguments do not need to be specified. Non-replacement means that once a treated case is matched to a non-treated case, both cases are removed from the pool.

To avoid dissimilar matches, we can constrain matches so that the absolute distance of propensity scores between two participants is less than a specified tolerance for matching or a caliper. The width of the caliper is by default in standard deviation units and can be specified using the `caliper` argument. A wide caliper may result in more matches and a larger sample, but inexact matching may occur as indicated by large distances on the propensity score between the treated and nontreated cases. Using varying caliper sizes can test the sensitivity of the findings. Here we use a caliper size of a quarter of a standard deviation, which is suggested by Rosenbaum and Rubin (1985).

The order of the matching can be specified using the `m.order` argument. If this argument is set to `largest`, then matching begins with the treated subject with the highest propensity score; if set to `smallest`, then matching takes places in ascending order of the distance measures; and if `random`, matching takes place in a random order. "The default is to go in descending order from the highest propensity score; doing so allows the units that would have the hardest time finding close matches to be matched first."[4]

Finally, the logit of the predicted probability from a logistic regression model can be supplied to the `distance` argument. The logit of the predicted probability is used, because the logit is approximately normally distributed.

The choice of explanatory variables (i.e., conditioning variables) in the first-step model predicting propensity scores of service receipt serves a paramount important role in the whole propensity score analysis. We chose these variables based on a review of substance abuse literature to determine what characteristics were associated with treatment receipt:

```r
# Logistic regression specification
(gm_f <- cobalt::f.build("aodserv", select(gm_df, PSH17A:other, -aodserv)))
```

```
## aodserv ~ PSH17A + CRA47A + married + high + bahigh + poverty2 +
##     poverty3 + poverty4 + poverty5 + employ + open + black +
##     hispanic + natam + cgrage1 + cgrage2 + cgrage3 + chdage1 +
##     chdage2 + chdage3 + mental + arrest + sexual + provide +
##     supervis + other
## <environment: 0x000001e8cc1b21d8>
```

We then calculate the logit of the predicted probability as the propensity score:

```r
gm_psm <- glm(gm_f, data = gm_df, family = binomial)
gm_ps <- predict(gm_psm, newdata = gm_df, type = "response")
gm_ps_logit <- log((1 - gm_ps) / gm_ps)
```

And then perform greedy nearest neighbor matching without replacement

```r
set.seed(1000)
(gm_out <- MatchIt::matchit(
  gm_f,
  data = gm_df,
  distance = gm_ps_logit,
  m.order = "largest", # descending order
  caliper = .25
))
```

```
## A matchit object
```

[4] https://kosukeimai.github.io/MatchIt/articles/matching-methods.html

```
##  - method: 1:1 nearest neighbor matching without replacement
##  - distance: User-defined [caliper]
##  - caliper: <distance> (0.311)
##  - number of obs.: 2758 (original), 566 (matched)
##  - target estimand: ATT
##  - covariates: PSH17A, CRA47A, married, high, bahigh, poverty2, poverty3, poverty4, poverty5, employ
```

Notice that a limitation of this matching scheme is that it reduces the sample size from 2758 to 566—283 cases in the control group and 283 cases in the treated group.

The `matchit` object will return a `match.matrix`, which contains the treated units as the rownames and the values in each row the names or indices of the control units matched to the treated units:

```
head(gm_out$match.matrix)
```

```
##    [,1]
## 2  "420"
## 3  "1852"
## 10 "2130"
## 14 NA
## 41 "782"
## 49 "1162"
```

Remember to use weights when estimating the treatment effect (but not necessary when performing 1:1 matching without replacement):

```
head(gm_out$weights)
```

```
## 1 2 3 4 5 6
## 0 1 1 0 0 0
```

1.1.5.1 Check Common Support Greedy matching is criticized, because it requires a sizeable common-support region to work. The common support region is defined as the region bounded by the maximum value of estimated propensity scores for the treated participants and by the minimum value of the estimated propensity scores for the nontreated participants. In this example, a sizeable common-support region exists. The `discard` argument in `matchit()` can be used to discard units outside a region of common support.

```
cobalt::bal.plot(gm_out, var.name = "distance")
```

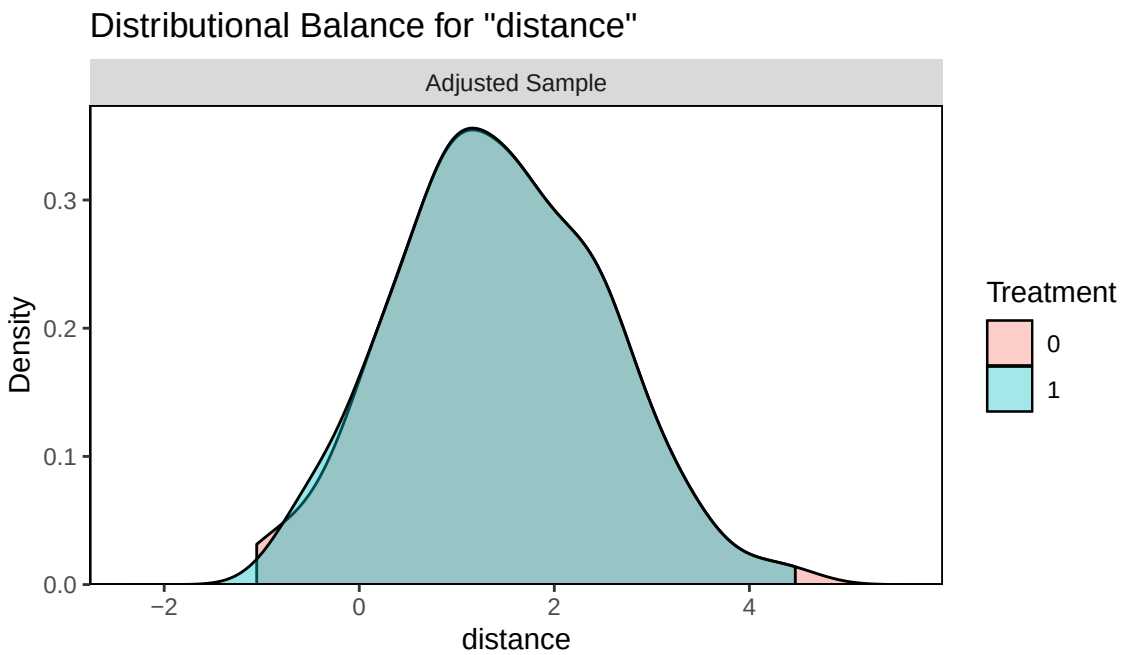

1.1.5.2 Check Balance

Covariate balance can be assessed using hypothesis tests, such as `chisq.test`:

```
gm_out_data <- MatchIt::match.data(gm_out)
chisq.test(gm_out_data$ra, gm_out_data$aodserv)
```

```
##
##  Pearson's Chi-squared test with Yates' continuity correction
##
## data:  gm_out_data$ra and gm_out_data$aodserv
## X-squared = 108.05, df = 1, p-value < 0.00000000000000022
```

The object from `matchit()` can be directly used in `cobalt` functions to produce balance tables and plots. To specify additional variables for which to display balance, use the argument `addl` in conjunction with `data`.

```
cobalt::love.plot(gm_out,
  binary = "std", threshold = c(m = .1), drop.distance = T,
  addl = c("ra", "cidi", "cgneed", "cwwrep"), data = gm_df
)
```

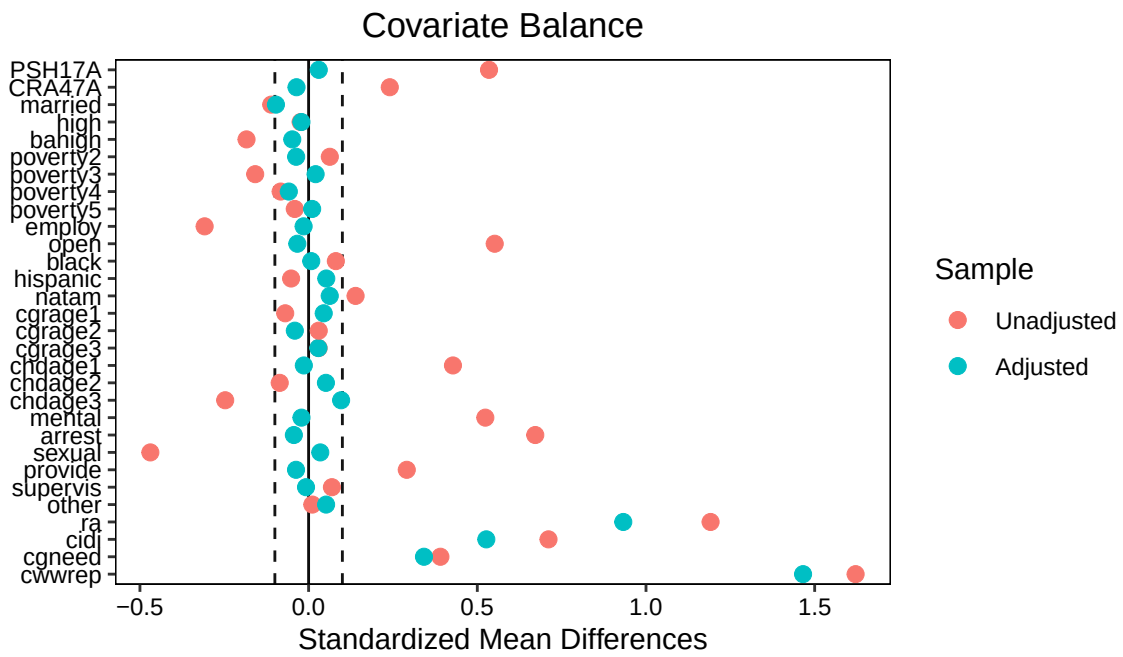


```
cobalt::bal.tab(gm_out,
  binary = "std", threshold = c(m = .1), un = T,
  addl = c("ra", "cidi", "cgneed", "cwwrep"), data = gm_df
)
```

```
## Call
##  MatchIt::matchit(formula = gm_f, data = gm_df, distance = gm_ps_logit,
##     m.order = "largest", caliper = 0.25)
##
## Balance Measures
##               Type Diff.Un Diff.Adj        M.Threshold
## distance Distance -1.3021   0.0068     Balanced, <0.1
## PSH17A     Binary  0.5347   0.0301     Balanced, <0.1
## CRA47A     Binary  0.2406  -0.0359     Balanced, <0.1
## married    Binary -0.1108  -0.0969     Balanced, <0.1
## high       Binary -0.0236  -0.0214     Balanced, <0.1
## bahigh     Binary -0.1840  -0.0485     Balanced, <0.1
## poverty2   Binary  0.0628  -0.0370     Balanced, <0.1
## poverty3   Binary -0.1584   0.0207     Balanced, <0.1
## poverty4   Binary -0.0829  -0.0587     Balanced, <0.1
## poverty5   Binary -0.0410   0.0105     Balanced, <0.1
## employ     Binary -0.3082  -0.0148     Balanced, <0.1
## open       Binary  0.5516  -0.0335     Balanced, <0.1
## black      Binary  0.0807   0.0078     Balanced, <0.1
## hispanic   Binary -0.0519   0.0524     Balanced, <0.1
## natam      Binary  0.1392   0.0626     Balanced, <0.1
## cgrage1    Binary -0.0694   0.0448     Balanced, <0.1
## cgrage2    Binary  0.0300  -0.0409     Balanced, <0.1
## cgrage3    Binary  0.0297   0.0289     Balanced, <0.1
## chdage1    Binary  0.4279  -0.0142     Balanced, <0.1
## chdage2    Binary -0.0856   0.0513     Balanced, <0.1
## chdage3    Binary -0.2473   0.0963     Balanced, <0.1
## mental     Binary  0.5238  -0.0212     Balanced, <0.1
```

```
## arrest     Binary  0.6718  -0.0438     Balanced, <0.1
## sexual     Binary -0.4690   0.0346     Balanced, <0.1
## provide    Binary  0.2909  -0.0373     Balanced, <0.1
## supervis   Binary  0.0690  -0.0077     Balanced, <0.1
## other      Binary  0.0110   0.0518     Balanced, <0.1
## ra         Binary  1.1915   0.9327 Not Balanced, >0.1
## cidi       Binary  0.7110   0.5267 Not Balanced, >0.1
## cgneed     Binary  0.3907   0.3419 Not Balanced, >0.1
## cwwrep     Binary  1.6213   1.4656 Not Balanced, >0.1
##
## Balance tally for mean differences
##                    count
## Balanced, <0.1        27
## Not Balanced, >0.1     4
##
## Variable with the greatest mean difference
##  Variable Diff.Adj        M.Threshold
##    cwwrep   1.4656 Not Balanced, >0.1
##
## Sample sizes
##           Control Treated
## All          2460     298
## Matched       283     283
## Unmatched    2177      15
```

1.1.6 Greedy Nearest Neighbor Mahalanobis Distance Matching Without Replacement

Here we perform Mahalanobis distance matching without replacement and without including an estimated propensity score.

```
set.seed(1000)
(gm_out2 <- MatchIt::matchit(
  gm_f,
  data = gm_df,
  method = "nearest",
  distance = "mahalanobis",
  replace = F
))
```

```
## A matchit object
##  - method: 1:1 nearest neighbor matching without replacement
##  - distance: Mahalanobis
##  - number of obs.: 2758 (original), 596 (matched)
##  - target estimand: ATT
##  - covariates: PSH17A, CRA47A, married, high, bahigh, poverty2, poverty3, poverty4, poverty5, employ
```

```
cobalt::love.plot(gm_out2,
  binary = "std", threshold = c(m = .1),
  addl = c("ra", "cidi", "cgneed", "cwwrep"), data = gm_df
)
```

1.1.6.1 Check Balance

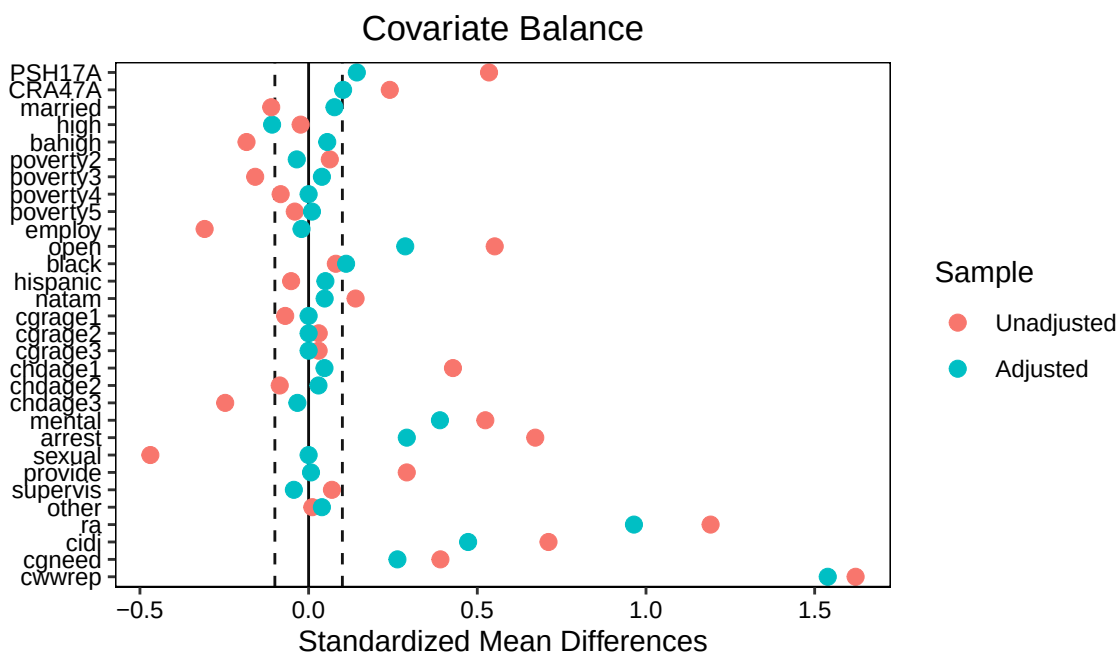


```
cobalt::bal.tab(gm_out2,
  binary = "std", threshold = c(m = .1),
  addl = c("ra", "cidi", "cgneed", "cwwrep"), data = gm_df
)
```

```
## Call
```

```
##  MatchIt::matchit(formula = gm_f, data = gm_df, method = "nearest",
##     distance = "mahalanobis", replace = F)
##
## Balance Measures
##            Type Diff.Adj        M.Threshold
## PSH17A   Binary   0.1429 Not Balanced, >0.1
## CRA47A   Binary   0.1022 Not Balanced, >0.1
## married  Binary   0.0767     Balanced, <0.1
## high     Binary  -0.1083 Not Balanced, >0.1
## bahigh   Binary   0.0552     Balanced, <0.1
## poverty2 Binary  -0.0351     Balanced, <0.1
## poverty3 Binary   0.0394     Balanced, <0.1
## poverty4 Binary   0.0000     Balanced, <0.1
## poverty5 Binary   0.0099     Balanced, <0.1
## employ   Binary  -0.0211     Balanced, <0.1
## open     Binary   0.2864 Not Balanced, >0.1
## black    Binary   0.1111 Not Balanced, >0.1
## hispanic Binary   0.0497     Balanced, <0.1
## natam    Binary   0.0476     Balanced, <0.1
## cgrage1  Binary   0.0000     Balanced, <0.1
## cgrage2  Binary   0.0000     Balanced, <0.1
## cgrage3  Binary   0.0000     Balanced, <0.1
## chdage1  Binary   0.0471     Balanced, <0.1
## chdage2  Binary   0.0292     Balanced, <0.1
## chdage3  Binary  -0.0333     Balanced, <0.1
## mental   Binary   0.3893 Not Balanced, >0.1
## arrest   Binary   0.2910 Not Balanced, >0.1
## sexual   Binary   0.0000     Balanced, <0.1
## provide  Binary   0.0071     Balanced, <0.1
## supervis Binary  -0.0437     Balanced, <0.1
## other    Binary   0.0394     Balanced, <0.1
## ra       Binary   0.9643 Not Balanced, >0.1
## cidi     Binary   0.4728 Not Balanced, >0.1
## cgneed   Binary   0.2633 Not Balanced, >0.1
## cwwrep   Binary   1.5388 Not Balanced, >0.1
##
## Balance tally for mean differences
##                     count
## Balanced, <0.1         19
## Not Balanced, >0.1     11
##
## Variable with the greatest mean difference
##  Variable Diff.Adj        M.Threshold
##    cwwrep   1.5388 Not Balanced, >0.1
##
## Sample sizes
##           Control Treated
## All          2460     298
## Matched       298     298
## Unmatched    2162       0
```

As seen above, balance has not been achieved in multiple covariates. According to Stuart (2010), "the Mahalanobis distance can work quite well when there are relatively few covariates (fewer than 8), but it does

not perform as well when the covariates are not normally distributed or there are many covariates."[5]

[5]Stuart, E. A. (2010). Matching Methods for Causal Inference: A Review and a Look Forward. *Statistical Science*, *25*(1), 1–21. https://doi.org/10.1214/09-STS313

1.2 Propensity Score Estimation Using Generalized Boosted Regression

1.2.1 Load Package

Generalized boosted regression (GBR) requires the `gbm` package.

```
library(gbm)
```

1.2.2 Load Data and Sort

Generalized boosted regression is an iterative method for creating propensity scores. Therefore, to create reproducible results, we need to use the `set.seed()` function.

After importing the data, missing data is deleted listwise, and the data is sorted randomly. According to the `gbm` package vignette, if the data is sorted in a systematic way, then the data should be shuffled before running `gbm`.

```
set.seed(1000)
gbr_df <- read_dta("data/g3aca1.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble() %>%
  select(
    intbl, ageyc, fmale, blck, whit, hisp, pcedu, ipovl, pcemft, fthr,
    dicsagg2, dicsint2, dccereg2, dccscom2, dccpros2, draggr2
  ) %>%
  drop_na() %>%
  add_column(runif = runif(nrow(.))) %>%
  arrange(runif) %>%
  select(-runif)
```

1.2.3 Fit Generalized Boosted Regression Model

The `gbm::gbm()` function has many arguments that can be fine-tuned. See `?gbm` for a detailed description of each argument. For example, to reduce prediction error, the `train.fraction` argument can be set to use a subsample of the observations for the estimation process. Here we also use `interaction.depth = 4` to specify a maximum of four splits for each simple tree used in the model, which allows all four-way interactions between all covariates to be considered for optimizing the likelihood function at each iteration. The `shrinkage` argument is also known as the learning rate or step-size reduction; we use a value of .0005 to ensure a smooth fit.

A summary of the fitted model provides us with *relative influence*, which is the percentage of log likelihood explained by each input variable. The percentages of influence for all predictor variables sum to 100%.

The GBM ouput showed that `blck` had the strongest influence on the likelihood function (33.7%), followed by `ageyc` (16.3%) and `draggr2` (9.4%).

```
(gbr_f <- cobalt::f.build("intbl", select(gbr_df, -intbl)))
```

```
## intbl ~ ageyc + fmale + blck + whit + hisp + pcedu + ipovl +
##     pcemft + fthr + dicsagg2 + dicsint2 + dccereg2 + dccscom2 +
##     dccpros2 + draggr2
## <environment: 0x000001e8c9536460>
```

```
set.seed(1000)
gbr_m1 <- gbm::gbm(
  formula = gbr_f,
  data = gbr_df,
  distribution = "bernoulli",
```

```
  n.trees = 1000, # number of trees to fit
  train.fraction = 0.8, # a random 80% subsample for estimation
  interaction.depth = 4, # allow all four-way interactions
  shrinkage = 0.0005 # small shrinkage to ensure smooth fit
)
summary(gbr_m1)
```

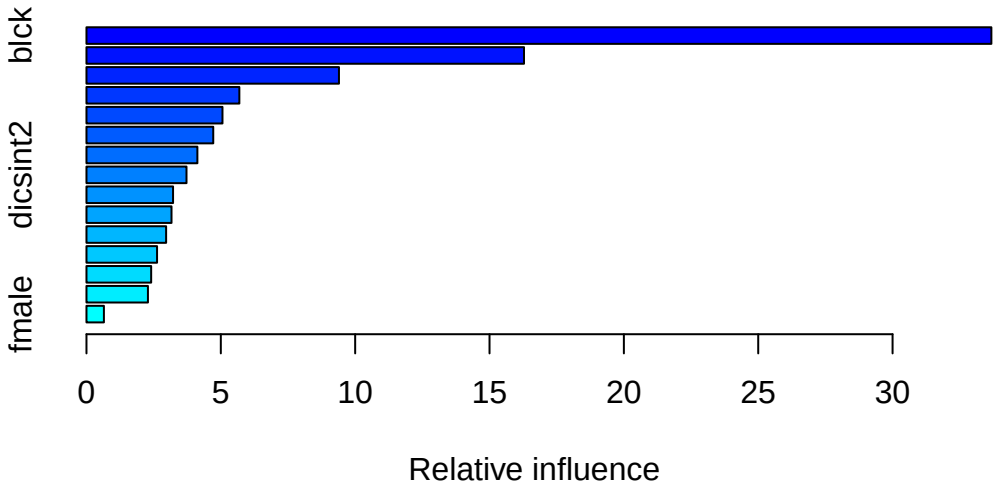


```
##                var    rel.inf
## blck          blck 33.6768868
## ageyc        ageyc 16.2834103
## draggr2   draggr2  9.3962151
## whit          whit  5.6895119
## ipovl        ipovl  5.0614691
## pcemft      pcemft  4.7179839
## pcedu        pcedu  4.1280063
## dicsint2 dicsint2  3.7223185
## dicsagg2 dicsagg2  3.2228567
## dccscom2 dccscom2  3.1648839
## dccereg2 dccereg2  2.9647178
## dccpros2 dccpros2  2.6269129
## hisp          hisp  2.4079157
## fthr          fthr  2.2874508
## fmale        fmale  0.6494603
```

1.2.4 Estimate Propensity Scores

After fitting the model, estimate propensity scores using the `predict.gbm()` function.

```
psb <- gbm::predict.gbm(gbr_m1, data = gbr_df, type = "response")
head(psb)
```

```
## [1] 0.4441531 0.4405950 0.5288983 0.5958331 0.3985049 0.5562076
```

1.2.5 Plot Propensity Score Distributions

As seen in the figure below, the propensity scores estimated by GBM has sufficient overlap between the control and treatment groups (i.e., "common support"), and the distributions are similar.

```
gbr_df %>%
  mutate(psb = psb, intbl = factor(intbl, labels = c("Control", "Treatment"))) %>%
  ggplot(aes(x = psb, color = intbl)) +
  theme_classic() +
  geom_density(size = 1) +
  xlim(0, 1) +
  ylim(0, 22) +
  labs(
    x = "Predicted Probability", y = "Density",
    title = "Propensity Scores Using Generalized Boosted Regression",
    color = "Treatment"
  ) +
  theme(legend.position = c(0.9, 0.85))
```

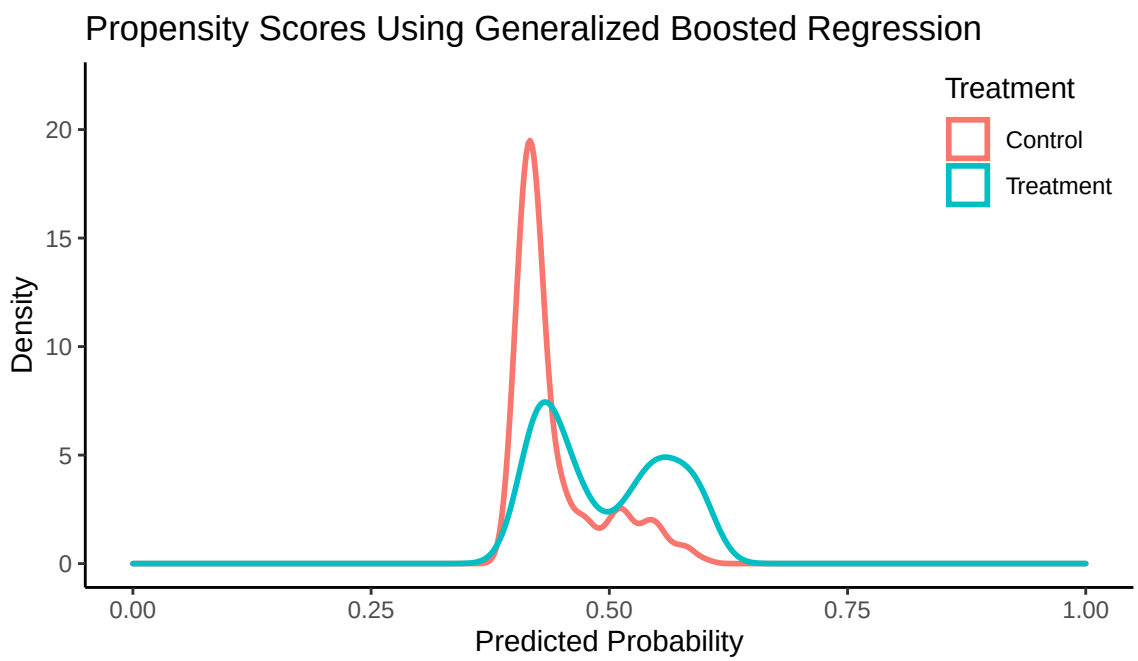

1.2.6 Summary Statistics of Propensity Scores

```
summary(psb)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.3913  0.4189  0.4383  0.4674  0.5202  0.6082
```

1.2.7 GBR Using the WeightIt Package

As an alternative to the `gbm` package, the `WeightIt` package can also fit GBR models.

```
set.seed(1000)
(gbr_m2 <- WeightIt::weightit(
  formula = gbr_f,
  data = gbr_df,
  method = "gbm",
  estimand = "ATE",
  distribution = "bernoulli",
  stop.method = "es.mean",
  n.trees = 10000, # different
  nTrain = 0.8 * nrow(gbr_df), # different
  interaction.depth = 4,
  shrinkage = 0.0005
))
```

```
## A weightit object
##  - method: "gbm" (propensity score weighting with GBM)
##  - number of obs.: 603
##  - sampling weights: none
##  - treatment: 2-category
##  - estimand: ATE
##  - covariates: ageyc, fmale, blck, whit, hisp, pcedu, ipovl, pcemft, fthr, dicsagg2, dicsint2, dccere
```

```
# Propensity Scores
head(gbr_m2$ps)
```

```
## [1] 0.3781517 0.6075811 0.6596871 0.7766516 0.2138153 0.7927817
```

```
# ATE Weights
head(gbr_m2$weights)
```

```
## [1] 2.644442 1.645871 1.515870 1.287579 1.271966 1.261381
```

1.2.7.1 Check Balance Using a standardized mean difference cut-off point of 0.1, it can be seen below that balance has been achieved in most, but not all, of the covariates:

```
cobalt::love.plot(gbr_m2,
  thresholds = c(m = .1),
  binary = "std", abs = T, drop.distance = T
) +
  labs(title = "Covariate Balance (ATE)")
```

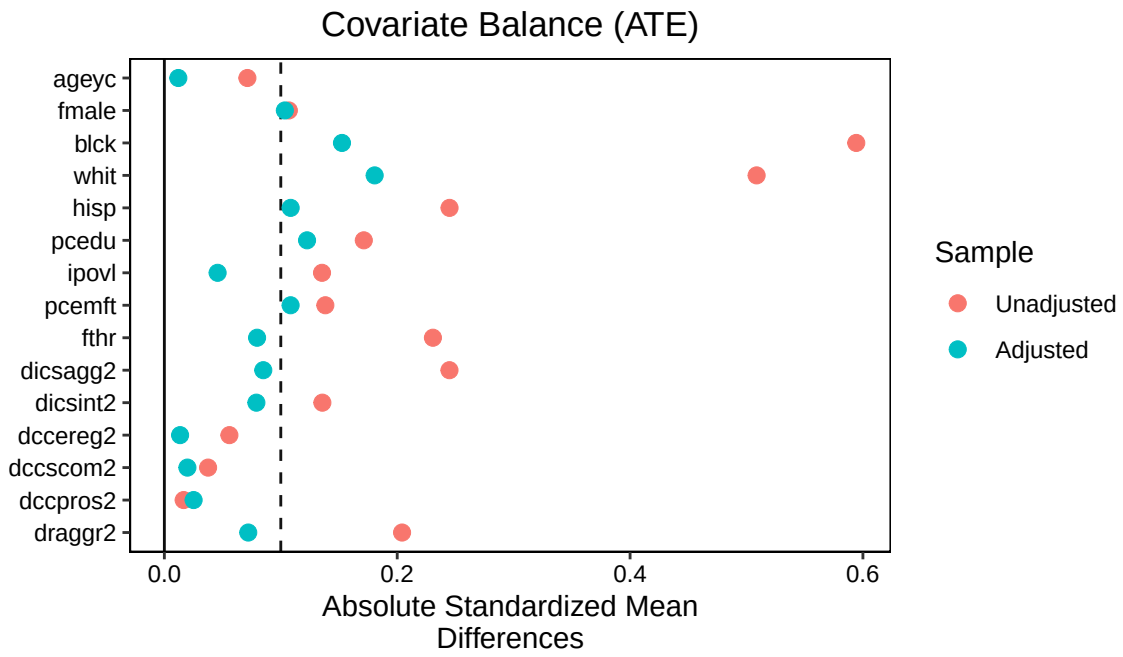

2 Day 2 - September 16, 2022

2.1 Optimal Matching

2.1.1 Load Packages

```
library(optmatch)
library(knitr)
library(broom)
library(sandwich)
library(rlang)
select <- dplyr::select
```

2.1.2 Load Data

This dataset comes from a study that investigates intergenerational dependence on welfare and its relation to child academic achievement.

The dependent variable is `lwss97`, the age-normed "letter-word" identification score of the Woodcock-Johnson Revised Tests of Achievement. A high score on this measure indicates high achievement. The treatment variable is `kuse` or children who ever used Aid to Families With Dependent Children (AFDC). The covariates are:

- `mratio96`: Ratio of Family Income to Poverty Line in 1996
- `pcged97`: Caregiver's Education in 1997 (Years of Schooling)
- `pcg_adc`: Caregiver's History of Using Welfare (Number of Years; range: 0-7)
- `black`: Child's Race: African American (1 = African American; 0 = Other)
- `age97`: Child's Age in 1997
- `male`: Child's Gender: Male (1 = Male; 0 = Female)

```
d <- haven::read_dta("data/opt/chpt5_2_original.dta")
df <- haven::read_dta("data/opt/chpt5_2.dta")
cds <- haven::read_dta("data/opt/chpt5_2ps.dta")
```

2.1.3 Bivariate Comparisons

With the exception of child's gender, the difference on each covariate between treated and control groups is statistically significant. Without controlling for these covariates, the study's estimate of treatment effect would be biased.

```
# Intergenerational dependence
gmodels::CrossTable(d$kuse, d$puse,
  prop.c = FALSE, prop.r = FALSE, prop.chisq = FALSE,
  chisq = TRUE
)
```

```
##
##
##    Cell Contents
## |-------------------------|
## |                       N |
## |         N / Table Total |
## |-------------------------|
##
##
## Total Observations in Table:  1003
##
```

```
##
##               | d$puse
##       d$kuse |         0 |         1 | Row Total |
## -------------|-----------|-----------|-----------|
##            0 |       615 |       114 |       729 |
##              |     0.613 |     0.114 |           |
## -------------|-----------|-----------|-----------|
##            1 |       154 |       120 |       274 |
##              |     0.154 |     0.120 |           |
## -------------|-----------|-----------|-----------|
## Column Total |       769 |       234 |      1003 |
## -------------|-----------|-----------|-----------|
##
##
## Statistics for All Table Factors
##
##
## Pearson's Chi-squared test
## ------------------------------------------------------------
## Chi^2 =  88.27405     d.f. =  1     p =  0.0000000000000000000005698621
##
## Pearson's Chi-squared test with Yates' continuity correction
## ------------------------------------------------------------
## Chi^2 =  86.70687     d.f. =  1     p =  0.0000000000000000000001258586
##
##
```

```r
# Wilcoxon Rank-Sum (Mann-Whitney) test and t-test
d %>%
  select(mratio96, pcged97, pcg_adc, black, age97, male, kuse) %>%
  pivot_longer(-kuse, names_to = "variable") %>%
  group_by(variable) %>%
  nest() %>%
  mutate(wilcoxon = map(data, ~ wilcox.test(.$value ~ .$kuse, correct = F))) %>%
  mutate(wilcoxon.stat = map(wilcoxon, ~ round(.$statistic, 3))) %>%
  mutate(wilcoxon.pvalue = map(wilcoxon, ~ round(.$p.value, 3))) %>%
  unnest(cols = c(wilcoxon.stat, wilcoxon.pvalue)) %>%
  mutate(ttest = map(data, ~ t.test(.$value ~ .$kuse, var.equal = T))) %>%
  mutate(t.stat = map(ttest, ~ round(.$statistic, 3))) %>%
  mutate(t.pvalue = map(ttest, ~ round(.$p.value, 3))) %>%
  unnest(cols = c(t.stat, t.pvalue)) %>%
  select(-data, -wilcoxon, -ttest)
```

```
## # A tibble: 6 x 5
## # Groups:   variable [6]
##   variable wilcoxon.stat wilcoxon.pvalue  t.stat t.pvalue
##   <chr>            <dbl>           <dbl>   <dbl>    <dbl>
## 1 mratio96       172618.           0      12.3      0
## 2 pcged97        147410.           0      12.3      0
## 3 pcg_adc         69728            0     -11.0      0
## 4 black           58038            0     -12.7      0
## 5 age97           87129            0.002  -3.11     0.002
## 6 male           101544            0.636   0.473    0.637
```

```r
# SMD Balance Checks
(c1 <- cobalt::bal.tab(
  x = d %>% select(mratio96, pcged97, pcg_adc, black, age97, male),
  treat = d$kuse,
  binary = "std",
  threshold = c(m = .1),
  s.d.denom = "pooled"
))
```

```
## Balance Measures
##             Type Diff.Un     M.Threshold.Un
## mratio96 Contin. -1.0158 Not Balanced, >0.1
## pcged97  Contin. -0.9067 Not Balanced, >0.1
## pcg_adc  Contin.  0.6703 Not Balanced, >0.1
## black     Binary  0.9343 Not Balanced, >0.1
## age97    Contin.  0.2196 Not Balanced, >0.1
## male      Binary -0.0335     Balanced, <0.1
##
## Balance tally for mean differences
##                    count
## Balanced, <0.1         1
## Not Balanced, >0.1     5
##
## Variable with the greatest mean difference
##  Variable Diff.Un     M.Threshold.Un
##  mratio96 -1.0158 Not Balanced, >0.1
##
## Sample sizes
##     Control Treated
## All     729     274
```

```r
cobalt::love.plot(c1)
```

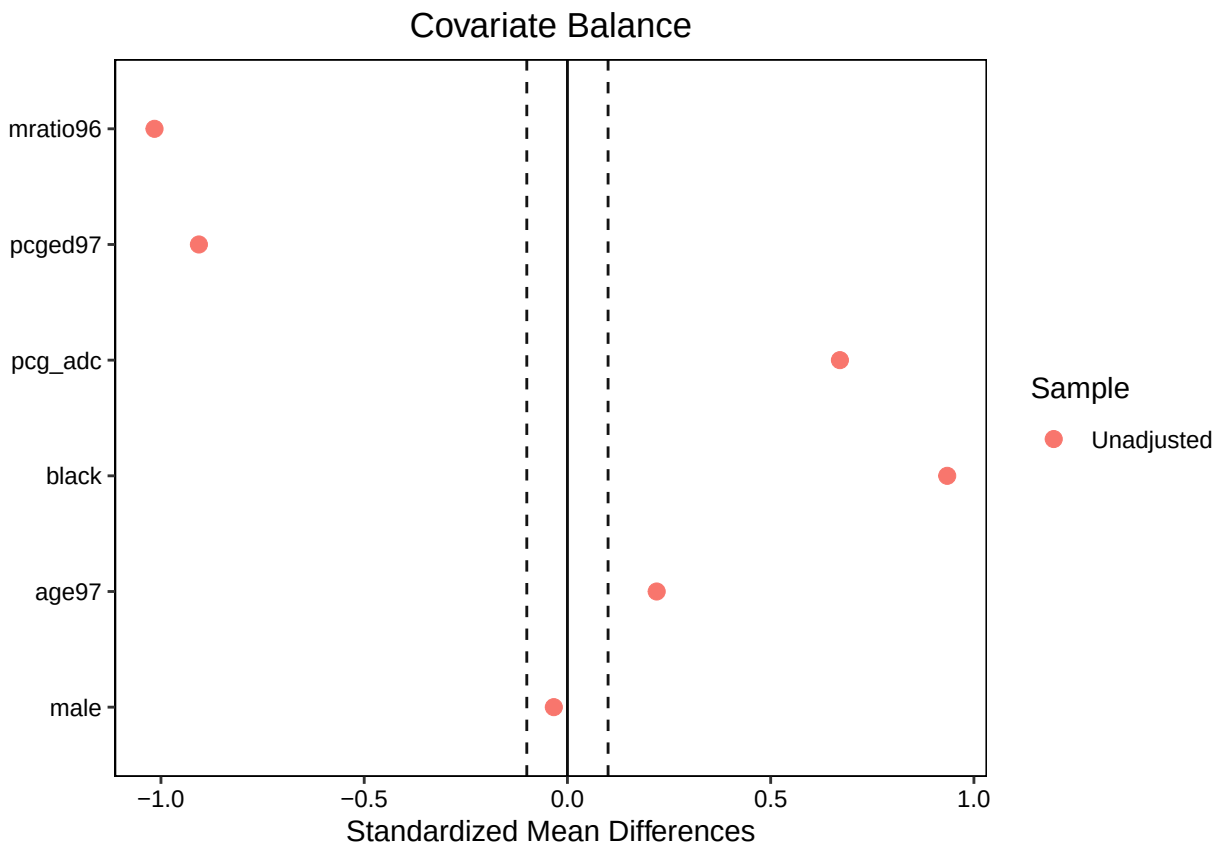

2.1.4 Outcome Models Without Bias Control

We first observe that the control group had a higher letter-word identification score than the treatment group:

```
d %>%
  ggplot(aes(y = lwss97, x = factor(kuse))) +
  geom_boxplot() +
  labs(x = "Treatment", y = "Letter-Word Identification Score") +
  scale_x_discrete(labels = c("Control", "Treated"))
```

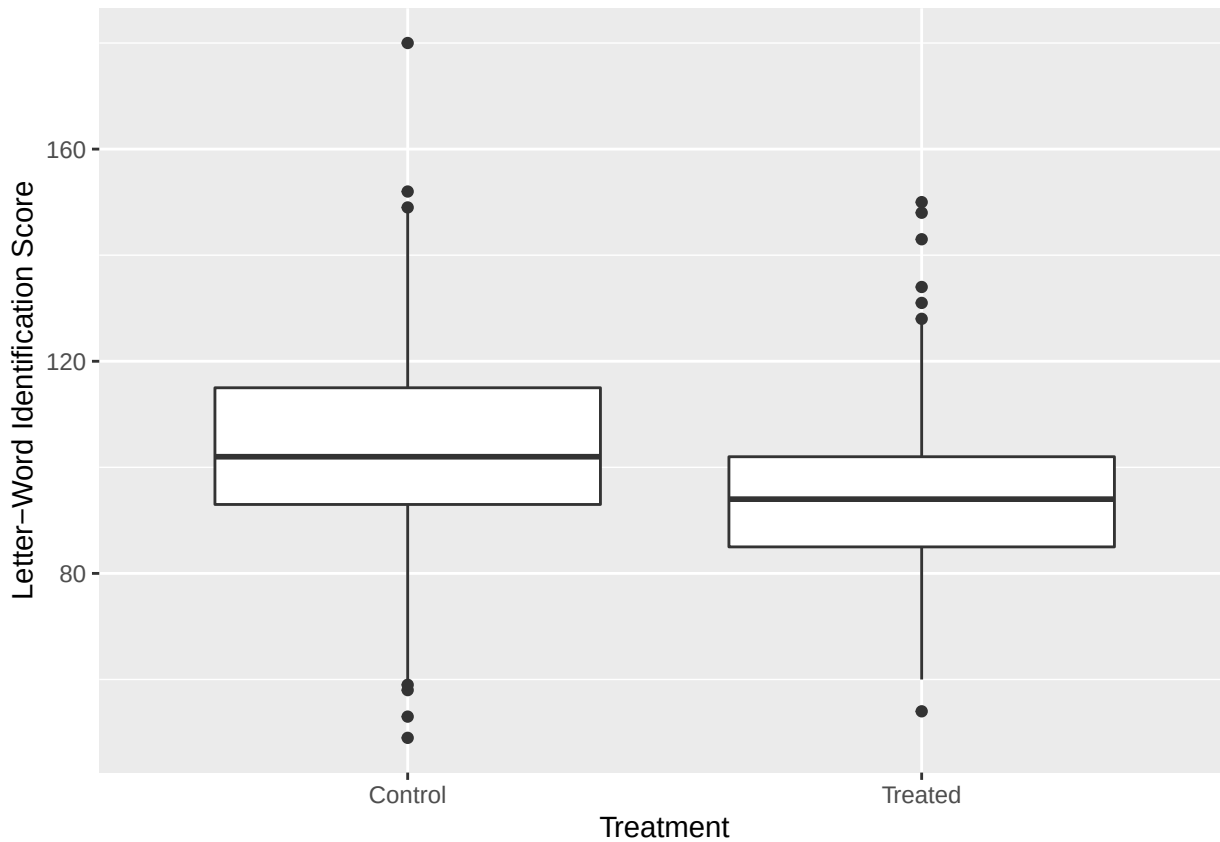

We can estimate the ATE using an independent samples t test. It shows that the treated group on average had a mean letter-word identification score that was 9.82 points lower than that of the control group ($p < .001$).

```
t.test(lwss97 ~ kuse, data = d, var.equal = T) %>%
  broom::tidy()
```

```
## # A tibble: 1 x 10
##   estimate estimate1 estimate2 statistic  p.value param~1 conf.~2 conf.~3 method
##      <dbl>     <dbl>     <dbl>     <dbl>    <dbl>   <dbl>   <dbl>   <dbl> <chr>
## 1     9.82      104.      94.2      8.51 6.27e-17    1001    7.56    12.1 Two S~
## # ... with 1 more variable: alternative <chr>, and abbreviated variable names
## #   1: parameter, 2: conf.low, 3: conf.high
```

We can also estimate the ATE using an OLS regression model. It shows that, controlling for covariates, the treated group on average had a letter-word identification score that is 4.73 points lower than that of the control group (p < .001).

```
reg <- lm(lwss97 ~ kuse + male + black + age97 + pcged97 + mratio96 + pcg_adc,
  data = d
)
broom::tidy(lmtest::coeftest(reg, vcov = sandwich::vcovCL, cluster = d$pcg_id))
```

```
## # A tibble: 8 x 5
##   term        estimate std.error statistic  p.value
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)   84.9       4.50      18.9  4.30e-68
## 2 kuse          -4.73      1.39      -3.40 7.00e- 4
## 3 male          -2.00      0.987     -2.02 4.34e- 2
## 4 black         -1.88      1.23      -1.53 1.27e- 1
```

```
## 5 age97          0.873     0.170      5.14 3.30e- 7
## 6 pcged97        0.910     0.336      2.71 6.91e- 3
## 7 mratio96       1.13      0.332      3.41 6.67e- 4
## 8 pcg_adc       -0.758     0.311     -2.44 1.50e- 2
```

2.1.5 Matching

Note: To be consistent with the STATA result, we will use the propensity scores created by Stata.

2.1.5.1 Distribution of estimated propensity scores As seen in the figures below, there is no sizeable common support region of the estimated propensity scores by treatment status, so using greedy matching would produce a nontrivial loss of matched participants. Therefore, we decide to use optimal matching.

```
cds %>%
  ggplot(aes(x = factor(kuse), y = ps)) +
  geom_boxplot() +
  labs(y = "Propensity Score", x = "Treatment Condition") +
  scale_x_discrete(labels = c("Control", "Treatment")) +
  theme_classic()
```

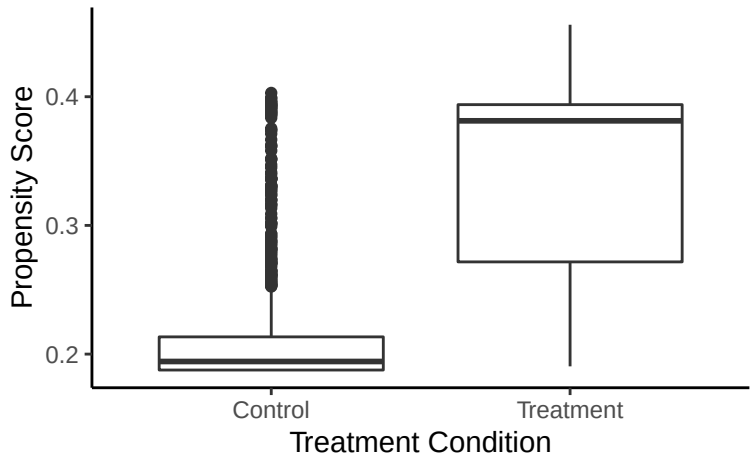


```
cds %>%
  ggplot(aes(x = ps, color = factor(kuse))) +
  geom_histogram(aes(y = ..density..), fill = "white", position = "dodge") +
  geom_density(alpha = .2, fill = "#FF6666") +
  labs(y = "Density", x = "Propensity Score", color = "Treatment Condition") +
  scale_color_discrete(labels = c("Control", "Treatment")) +
  theme_classic() +
  theme(legend.position = "bottom")
```

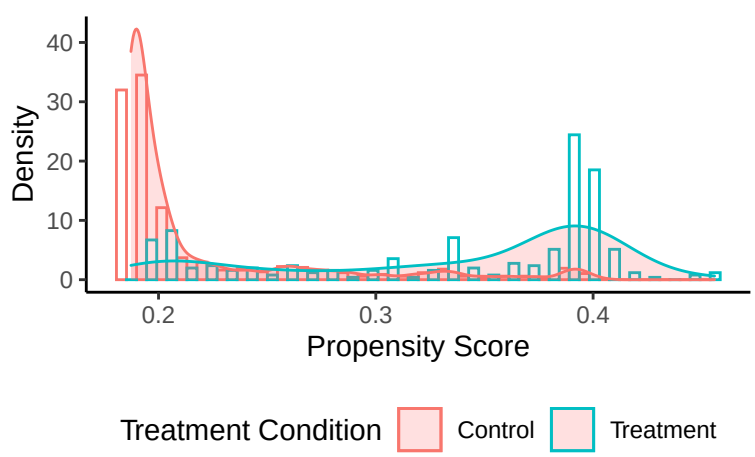

We would lose 45.3% cases if we used greedy nearest neighbor matching:

```r
(matchit(kuse ~ mratio96 + pcged97 + pcg_adc + black + age97 + male,
  data = d,
  method = "nearest",
  distance = "glm"
))
```

```
## A matchit object
##  - method: 1:1 nearest neighbor matching without replacement
##  - distance: Propensity score
##              - estimated with logistic regression
##  - number of obs.: 1003 (original), 548 (matched)
##  - target estimand: ATT
##  - covariates: mratio96, pcged97, pcg_adc, black, age97, male
```

```r
(nrow(d) - 548) / nrow(d)
```

```
## [1] 0.4536391
```

2.1.5.2 Optimal Matching The first argument that `optmatch::fullmatch()` requires is "a matrix of non-negative discrepancies, each indicating the permissibility and desirability of matching the unit corresponding to its row (treated) to the unit corresponding to its column (control)." First, we calculate the sample ranks of the propensity scores generated by generalized boosted regression using `rank(ps)`. Second, we use `outer(x, y, "-")`, which will subtract y control units from x treated units, to form a matrix with rows equal to the number of treated subjects and columns equal to the number of control subjects. Finally, we take the absolute value of these differences.

By default, `min.controls` or the minimum ratio of controls to treatments that is to be permitted within a matched set is zero, and `max.controls` or the maximum ratio of controls to treatments is infinite.

Hansen found that in the context of a specific application, variable matching with a specific structure worked best; that is, each treated participant was matched to at least $\frac{.5(1-\hat{P})}{\hat{P}}$ controls and at most $\frac{2(1-\hat{P})}{\hat{P}}$ controls, where $\hat{P}$ represents the proportion of treated participants in the sample. In this example, $\hat{P} = \frac{274}{729+274} = 27.32\%$, so the number of minimum controls is $\frac{.5(1-.2732)}{.2732} = 1.33$, and the number of maximum controls is $\frac{2(1-.2732)}{.2732} = 5.32$.

For `optmatch::pairmatch()`, we can set the number of controls to be matched to each treatment using the `controls` argument.

```
attach(cds)
prank <- rank(ps)
names(prank) <- kid
d1 <- outer(prank[kuse == 1], prank[kuse == 0], "-")
d1 <- abs(d1)
table(cds$kuse) # 729 control, 274 treated
dim(d1) # 274 x 729

# Full matching
fm <- fullmatch(d1)
(fm.d <- matched.distances(fm, d1, pres = TRUE))
unlist(fm.d, max)

# Variable Matching (at least 1, at most 4)
vm1 <- fullmatch(d1, min.controls = 1, max.controls = 4)
(vm1.d <- matched.distances(vm1, d1, pres = TRUE))
unlist(vm1.d, max)

# Variable Matching (at least 2, at most 4)
vm2 <- fullmatch(d1, min.controls = 2, max.controls = 4)
(vm2.d <- matched.distances(vm2, d1, pres = TRUE))
unlist(vm2.d, max)

# Variable Matching (using Hensen's equation)
vm3 <- fullmatch(d1, min.controls = 1.33, max.controls = 5.32)
(vm3.d <- matched.distances(vm3, d1, pres = TRUE))
unlist(vm3.d, max)

# Variable Matching (at least 2, at most 7)
vm4 <- fullmatch(d1, min.controls = 2, max.controls = 7)
(vm4.d <- matched.distances(vm4, d1, pres = TRUE))
unlist(vm4.d, max)

# 1:1 Pair Matching
pm <- pairmatch(d1, controls = 1)
```

```r
(pm.d <- matched.distances(pm, d1, pres = TRUE))
unlist(pm.d, max)
```

Optimal matching aims to minimize the total sample distance of propensity scores. We can extract the distances of matched units from their matched counterpairs using `optmatch::matched.distances()` and then calculate the total distance using `sum(unlist())`. The ratio of treatment to control and the number of strata for each structure can be inspected with `optmatch::stratumStructure()`.

```r
mean(unlist(fm.d))
```

```
## [1] 35.89389
```

```r
sum(unlist(fm.d))
```

```
## [1] 31120
```

```r
stratumStructure(fm)
```

```
##  25:1  14:1  11:1  10:1   9:1   7:1   6:1   5:1   4:1   3:1   2:1   1:1   1:2
##     1     1     2     2     1     1     2     3     3     5     8    48     8
##   1:3   1:4   1:5   1:6   1:7   1:8   1:9  1:10  1:11  1:15  1:17  1:18  1:21
##    10     7     8     3     2     4     1     5     2     1     1     2     1
##  1:24  1:31  1:64 1:185
##     1     1     1     1
```

```r
mean(unlist(vm1.d))
```

```
## [1] 328.7579
```

```r
sum(unlist(vm1.d))
```

```
## [1] 239664.5
```

```r
stratumStructure(vm1)
```

```
## 1:1 1:3 1:4
## 122   1 151
```

```r
mean(unlist(vm2.d))
```

```
## [1] 348.1852
```

```r
sum(unlist(vm2.d))
```

```
## [1] 253827
```

```r
stratumStructure(vm2)
```

```
## 1:2 1:3 1:4
## 183   1  90
```

```r
mean(unlist(vm3.d))
```

```
## [1] 269.7222
```

```r
sum(unlist(vm3.d))
```

```
## [1] 196627.5
```

```r
stratumStructure(vm3)
```

```
## 1:1 1:6
## 183  91
```

```r
mean(unlist(vm4.d))
```

```
## [1] 312.3162
```

```r
sum(unlist(vm4.d))
```

```
## [1] 227678.5
```

```r
stratumStructure(vm4)
```

```
## 1:2 1:3 1:7
## 237   1  36
```

```r
mean(unlist(pm.d))
```

```
## [1] 143.2263
```

```r
sum(unlist(pm.d))
```

```
## [1] 39244
```

```r
stratumStructure(pm)
```

```
## 1:1 0:1
## 274 455
```

```r
# Define covariates
cov_labels <- tibble(
  cov = c(
    "mratio96", "pcged97", "pcg_adc",
    "black", "age97", "male"
  ),
  new = c(
    "Ratio of family income to poverty line in 1996",
    "Caregiver's education in 1997 (years of schooling)",
    "Caregiver's number of years using AFDC in childhood",
    "Child's race: African American (reference: other)",
    "Child's age in 1997", "Child's gender: male (reference: female)"
  ),
  order = c(1, 2, 3, 4, 5, 6)
)

# Calculate dx and dxm using imbalance()
arg1 <- rep(c("mratio96", "pcged97", "pcg_adc", "black", "age97", "male"), each = 7)
arg2 <- rep(c("before", "fm", "vm1", "vm2", "vm3", "vm4", "pm"), length(arg1) / 7)
table_5.10 <- map2_dfr(arg1, arg2, imbalance) %>%
  mutate(method = recode(method,
    `before` = "Before Matching",
    `fm` = "Full Matching",
    `vm1` = "Variable Matching 1 (at least 1, at most 4)",
    `vm2` = "Variable Matching 2 (at least 2, at most 4)",
    `vm3` = "Variable Matching 3 (Hansen's equation)",
    `vm4` = "Variable Matching 4 (at least 2, at most 7)",
    `pm` = "Pair matching"
  )) %>%
  left_join(cov_labels, by = "cov") %>%
  arrange(order) %>%
```

```r
  select(cov, matching_scheme = method, dx, dxm)

# Print Table
print(table_5.10, n = 42)
```

2.1.5.3 Covariate Imbalance Before and After Matching by Matching Scheme

```
## # A tibble: 42 x 4
##    cov      matching_scheme                                  dx      dxm
##    <chr>    <chr>                                         <dbl>    <dbl>
##  1 mratio96 Before Matching                               1.02   NA
##  2 mratio96 Full Matching                                NA       0.0408
##  3 mratio96 Variable Matching 1 (at least 1, at most 4) NA       0.790
##  4 mratio96 Variable Matching 2 (at least 2, at most 4) NA       0.940
##  5 mratio96 Variable Matching 3 (Hansen's equation)     NA       0.736
##  6 mratio96 Variable Matching 4 (at least 2, at most 7) NA       0.869
##  7 mratio96 Pair matching                                NA       0.253
##  8 pcged97  Before Matching                               0.907  NA
##  9 pcged97  Full Matching                                NA       0.168
## 10 pcged97  Variable Matching 1 (at least 1, at most 4) NA       0.722
## 11 pcged97  Variable Matching 2 (at least 2, at most 4) NA       0.867
## 12 pcged97  Variable Matching 3 (Hansen's equation)     NA       0.752
## 13 pcged97  Variable Matching 4 (at least 2, at most 7) NA       0.814
## 14 pcged97  Pair matching                                NA       0.340
## 15 pcg_adc  Before Matching                               0.670  NA
## 16 pcg_adc  Full Matching                                NA       0.00737
## 17 pcg_adc  Variable Matching 1 (at least 1, at most 4) NA       0.560
## 18 pcg_adc  Variable Matching 2 (at least 2, at most 4) NA       0.649
## 19 pcg_adc  Variable Matching 3 (Hansen's equation)     NA       0.551
## 20 pcg_adc  Variable Matching 4 (at least 2, at most 7) NA       0.626
## 21 pcg_adc  Pair matching                                NA       0.403
## 22 black    Before Matching                               0.933  NA
## 23 black    Full Matching                                NA       0.0106
## 24 black    Variable Matching 1 (at least 1, at most 4) NA       0.788
## 25 black    Variable Matching 2 (at least 2, at most 4) NA       0.919
## 26 black    Variable Matching 3 (Hansen's equation)     NA       0.732
## 27 black    Variable Matching 4 (at least 2, at most 7) NA       0.848
## 28 black    Pair matching                                NA       0.442
## 29 age97    Before Matching                               0.220  NA
## 30 age97    Full Matching                                NA       0.0690
## 31 age97    Variable Matching 1 (at least 1, at most 4) NA       0.215
## 32 age97    Variable Matching 2 (at least 2, at most 4) NA       0.219
## 33 age97    Variable Matching 3 (Hansen's equation)     NA       0.190
## 34 age97    Variable Matching 4 (at least 2, at most 7) NA       0.238
## 35 age97    Pair matching                                NA       0.137
## 36 male     Before Matching                               0.0335 NA
## 37 male     Full Matching                                NA       0.0532
## 38 male     Variable Matching 1 (at least 1, at most 4) NA       0.0219
## 39 male     Variable Matching 2 (at least 2, at most 4) NA       0.0487
## 40 male     Variable Matching 3 (Hansen's equation)     NA       0.00974
## 41 male     Variable Matching 4 (at least 2, at most 7) NA       0.0325
## 42 male     Pair matching                                NA       0.0873
```

2.1.6 Post-Full-Matching Analysis of Outcome

2.1.6.1 Hodges-Lehmann Aligned Rank Test The Hodges-Lehmann aligned rank test can be used on matched samples created by full matching or variable matching. We used full matching and found that children who used AFDC had a letter-word identification score in 1997 that was on average 1.97 points lower than those who had never used AFDC; the difference was statistically significant at a .05 level (one-tailed). We used the Hodges-Lehmann test to gauge the statistical significance. The study also detected an effect size of .19, which is a small effect size in terms of Cohen's (1988) criteria.

```
hodgesl(df, lwss97, fm, kuse)
```

```
## # A tibble: 1 x 6
##   HL_mean HL_se     z      p     i tx_effect
##     <dbl> <dbl> <dbl>  <dbl> <int>     <dbl>
## 1  -7039. 3247. -2.17 0.0151     1     -1.97
```

```
imbalance2(df, lwss97, kuse, fm)$dxm # dxm for the outcome variable is Cohen's d
```

2.1.6.2 Cohen's d

```
## [1] 0.1896308
```

2.1.7 Post-Matching Analysis Using Regression of Difference Scores

A regression of difference scores may be performed on matched samples created by optimal pair matching. Based on the pair-matched sample, we first calculated difference scores between treated and control cases for each pair on all study variables (i.e., on the outcome variable and all covariates). We then regressed the difference score of the outcome on the difference scores of covariates. In addition, note that our model includes a correction for clustering effects (i.e., children are nested within caregivers) that we accomplished by using robust estimates of standard errors. The intercept of a difference score regression indicates the ATE of the sample. The estimated intercept from this model is -3.17 ($p < .05$). Thus, using pair matching and regression adjustment, the study found that, on average, children who used AFDC had a letter-word identification score in 1997 that was 3.17 points lower than do children who never used AFDC; this finding was statistically significant.

```
# kuse == 1
df1 <- df %>%
  select(pm, kuse,
    y1 = lwss97, male1 = male, black1 = black, age971 = age97,
    pcged971 = pcged97, mratio961 = mratio96, pcg_id
  ) %>%
  filter(kuse == 1)

# kuse == 0
df0 <- df %>%
  select(pm, kuse,
    y0 = lwss97, male0 = male, black0 = black, age970 = age97,
    pcged970 = pcged97, mratio960 = mratio96, pcg_id
  ) %>%
  filter(kuse == 0)

# Merge Data
df2 <- left_join(df0, df1, "pm")
df2 <- df2 %>%
  mutate(
    y = y1 - y0,
```

```r
    male = male1 - male0,
    black = black1 - black0,
    age97 = age971 - age970,
    pcged97 = pcged971 - pcged970,
    mratio96 = mratio961 - mratio960
  )

# Regression of Difference Scores (one-tailed test)
reg1 <- lm(y ~ male + black + age97 + pcged97 + mratio96, data = df2)
broom::tidy(lmtest::coeftest(reg1, vcov = vcovCL, cluster = df2$pcg_id.x)) %>%
  mutate(p.value.one.tailed = p.value / 2)
```

```
## # A tibble: 6 x 6
##   term        estimate std.error statistic p.value p.value.one.tailed
##   <chr>          <dbl>     <dbl>     <dbl>   <dbl>              <dbl>
## 1 (Intercept)  -3.17       1.73    -1.83    0.0682             0.0341
## 2 male         -1.21       1.84    -0.657   0.512              0.256
## 3 black        -3.34       2.33    -1.43    0.153              0.0765
## 4 age97         0.229      0.329    0.695   0.487              0.244
## 5 pcged97       0.0111     0.508    0.0219  0.983              0.491
## 6 mratio96      3.15       1.67     1.89    0.0602             0.0301
```

2.1.8 Optimal Matching Using the MatchIt Package

For an easier interface for the `optmatch` package, the `MatchIt::matchit()` function also implements optimal full matching and optimal pair matching with the `method = "full"` and `method = "optimal"` arguments, respectively. Variable matching can be set using the `ratio`, `min.controls`, and `max.controls` arguments.

Below is an illustration of optimal pair matching (by default, 1:1 without replacement) using generalized boosted model to generate the propensity scores and then a regression model for the outcome that incorporates the matching weights. Here we find that on average children who used AFDC had a letter-word identification score in 1997 that was 4.14 points lower than do children who never used AFDC ($p < .001$).

```r
# Optimal 1:1 Pair Matching Without Replacement Using GBM (ATT)
set.seed(1000)
(pm2 <- matchit(
  formula = kuse ~ pcg_adc + age97 + mratio96 + pcged97 + black,
  data = d,
  method = "optimal",
  replace = F,
  ratio = 1,
  distance = "gbm",
  estimand = "ATT"
))

# Regression with Matching Weights and Robust Standard Errors
pm2.lm <- lm(lwss97 ~ kuse + male + black + age97 + pcged97 + mratio96, data = d,
             weights = pm2$weights)
broom::tidy(lmtest::coeftest(pm2.lm, vcov = vcovCL, cluster = d$pcg_id))

# Balance Checks
cobalt::bal.tab(pm2, binary = "std", un = T, threshold = c(m = .1))
cobalt::love.plot(pm2, un = T, binary = "std", threshold = c(m = .1))
```

2.2 Propensity Score Weighting

2.2.1 Load Packages

Propensity score weighting can be accomplished with base R functions. However, we need the `lmtest` and `sandwich` packages to estimate clustered covariance matrices in this example. Using these packages, we can obtain estimates and standard errors that are identical to Stata's `regress` program.

```
library(lmtest)
library(sandwich)
```

2.2.2 Description of Dataset

This dataset is from a study that investigates intergenerational dependence on welfare and its relation to child academic achievement.[6]

The dependent variable is `lwss97` or "letter-word identification" score, and the treatment condition is `kuse` or children who used Aid to Families With Dependent Children (AFDC). The covariates are:

- `male`: Child's Gender: Male (Reference: Female)
- `black`: Child's Race: African American (Reference: Other)
- `age97`: Child's Age in 1997
- `pcged97`: Caregiver's Education in 1997 (Years of Schooling)
- `mratio96`: Ratio of Family Income to Poverty Line in 1996

Additionally, `pcg_id` is a cluster variable that identifies children nested within families.

2.2.3 Estimate ATE and ATT Weights

Separate weights need to be calculated for estimating the average treatment effect (ATE) and the average treatment effect for the treated (ATT). Propensity score weighting for estimating ATE is generally referred to as the inverse probability of treatment weights (IPTW) estimator.

For ATE, the weight estimates are calculated as follows for the treatment group:

$$\omega = \frac{1}{\hat{e}(x)}$$

And for the control group:

$$\omega = \frac{1}{1 - \hat{e}(x)}$$

For ATT, the weight is 1 for a treated case. The weight for a comparison case is:

$$\omega = \frac{\hat{e}(x)}{1 - \hat{e}(x)}$$

[6] Hofferth, S., Stafford, F. P., Yeung, W. J., Duncan, G. J., Hill, M. S., Lepkowski, J., et al. (2001). *Panel study of income dynamics, 1968–1999: Supplemental files (computer file), ICPSR version.* Ann Arbor: University of Michigan Survey Research Center.

2.2.4 Load Data with Propensity Scores and Calculate Weights

```
psw_df <- read_dta("data/chpt5_2_original.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble() %>%
  mutate(
    ate_w = ifelse(kuse == 0, 1 / (1 - ps), 1 / ps),
    att_w = ifelse(kuse == 0, ps / (1 - ps), 1)
  )
```

2.2.5 Calculate Weights with the WeightIt Package

```
# Load Package
library(WeightIt)
```

```
## Warning: package 'WeightIt' was built under R version 4.2.1
```

```
# Estimate ATE and ATT weights and Compare with Previous Results
ate_w2 <- WeightIt::get_w_from_ps(ps = psw_df$ps, treat = psw_df$kuse, estimand = "ATE")
table(ate_w2 == psw_df$ate_w)
```

```
##
## TRUE
## 1003
```

```
att_w2 <- WeightIt::get_w_from_ps(ps = psw_df$ps, treat = psw_df$kuse, estimand = "ATT")
table(att_w2 == (psw_df$ate_w * psw_df$ps))
```

```
##
## TRUE
## 1003
```

2.2.6 Outcome Analysis

2.2.6.1 Weighted Regression with ATE Weights After creating the weights, use the `weights` argument in `lm()` to run a weighted outcome analysis and `lmtest::coeftest()` to control for clustering effects.

This analysis showed that children who used Aid to Families With Dependent Children (AFDC) had an average letter-word identification score that was 5.16 points lower than children who never used AFDC, $p < .01$.

```
psw_ate <- lm(lwss97 ~ kuse + male + black + age97 + pcged97 + mratio96,
  data = psw_df, weights = ate_w
)
lmtest::coeftest(psw_ate, vcov. = vcovCL(psw_ate, cluster = psw_df$pcg_id))
```

```
##
## t test of coefficients:
##
##             Estimate Std. Error t value              Pr(>|t|)
## (Intercept) 84.19992    4.83032 17.4315 < 0.00000000000000022 ***
## kuse        -5.16399    1.42438 -3.6254             0.0003031 ***
## male        -1.62201    1.09186 -1.4855             0.1377180
## black       -2.49898    1.34670 -1.8556             0.0638009 .
## age97        0.73868    0.18075  4.0867            0.00004727 ***
```

```
## pcged97      0.99264    0.35596  2.7886             0.0053938 **
## mratio96     1.13856    0.32220  3.5337             0.0004286 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

2.2.6.2 Weighted Regression with ATT Weights When considering only individuals assigned to the treatment condition, children who used AFDC had an average letter-word identification score that was 4.62 points lower than children who never used AFDC, $p < .01$.

```
psw_att <- lm(lwss97 ~ kuse + male + black + age97 + pcged97 + mratio96,
  data = psw_df, weights = att_w
)
lmtest::coeftest(psw_att, vcov. = vcovCL(psw_att, cluster = psw_df$pcg_id))
```

```
##
## t test of coefficients:
##
##             Estimate Std. Error t value              Pr(>|t|)
## (Intercept) 85.29467    5.05331 16.8790 < 0.00000000000000022 ***
## kuse        -4.62058    1.41182 -3.2728              0.001102 **
## male        -1.58995    1.14705 -1.3861              0.166020
## black       -2.74605    1.41605 -1.9392              0.052756 .
## age97        0.61577    0.20001  3.0787              0.002136 **
## pcged97      0.92698    0.36718  2.5246              0.011738 *
## mratio96     1.26018    0.33556  3.7555              0.000183 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

2.2.7 Check Balance

To assess balance before and after propensity score weighting, use weighted logistic regression for dummy covariates and weighted simple regression for continuous covariates. Some examples are included below, and the full code can be found in Section 7.3.1 of the PSA-R code.

In model `psw_c3` below, the treatment dummy variable is significant, meaning that there is no sufficient balance after the propensity score weighting.

To assess balance before propensity score weighting, remove the `weights` argument.

```
psw_c1 <- glm(male ~ kuse, family = quasibinomial, data = psw_df, weights = ate_w)
lmtest::coeftest(psw_c1, vcov. = vcovCL(psw_c1, cluster = psw_df$pcg_id))
```

```
##
## z test of coefficients:
##
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.142911   0.075538  1.8919   0.0585 .
## kuse        -0.060271   0.150143 -0.4014   0.6881
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
psw_c2 <- glm(male ~ kuse, family = quasibinomial, data = psw_df, weights = att_w)
lmtest::coeftest(psw_c2, vcov. = vcovCL(psw_c2, cluster = psw_df$pcg_id))
```

```
##
## z test of coefficients:
##
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.152522   0.077410  1.9703   0.0488 *
## kuse        -0.079497   0.147480 -0.5390   0.5899
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
psw_c3 <- lm(age97 ~ kuse, weights = ate_w, data = psw_df)
lmtest::coeftest(psw_c3, vcov. = vcovCL(psw_c3, cluster = psw_df$pcg_id))
```

```
##
## t test of coefficients:
##
##             Estimate Std. Error t value              Pr(>|t|)
## (Intercept)  6.51389    0.10900 59.7615 < 0.00000000000000022 ***
## kuse         0.61064    0.21883  2.7905              0.005362 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
psw_c4 <- lm(age97 ~ kuse, weights = att_w, data = psw_df)
lmtest::coeftest(psw_c4, vcov. = vcovCL(psw_c4, cluster = psw_df$pcg_id))
```

```
##
## t test of coefficients:
##
##             Estimate Std. Error t value             Pr(>|t|)
## (Intercept)  6.55501    0.11499 57.0057 < 0.0000000000000002 ***
## kuse         0.56178    0.21939  2.5606              0.01059 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Alternatively, balance can be assessed using standardized mean differences:

```
# ATE
cobalt::bal.tab(
  x = select(psw_df, male, black, age97, pcged97, mratio96),
  treat = psw_df$kuse,
  weights = psw_df$ate_w,
  binary = "std",
  continuous = "std",
  s.d.denom = "pooled",
  un = T,
  stats = c("mean.diffs"),
  thresholds = c(m = .1)
)
```

```
## Balance Measures
##             Type Diff.Un Diff.Adj        M.Threshold
## male      Binary -0.0335  -0.0301     Balanced, <0.1
## black     Binary  0.9343   0.8006 Not Balanced, >0.1
## age97    Contin.  0.2196   0.2181 Not Balanced, >0.1
## pcged97  Contin. -0.9067  -0.8244 Not Balanced, >0.1
## mratio96 Contin. -1.0158  -0.8775 Not Balanced, >0.1
##
## Balance tally for mean differences
##                    count
## Balanced, <0.1         1
## Not Balanced, >0.1     4
```

```
##
## Variable with the greatest mean difference
##  Variable Diff.Adj        M.Threshold
##  mratio96  -0.8775 Not Balanced, >0.1
##
## Effective sample sizes
##            Control Treated
## Unadjusted  729.    274.
## Adjusted    725.43  257.71
```

```r
# ATT
cobalt::bal.tab(
  x = select(psw_df, male, black, age97, pcged97, mratio96),
  treat = psw_df$kuse,
  weights = psw_df$att_w,
  binary = "std",
  continuous = "std",
  s.d.denom = "treated",
  un = T,
  stats = c("mean.diffs"),
  thresholds = c(m = .1)
)
```

```
## Balance Measures
##             Type Diff.Un Diff.Adj        M.Threshold
## male      Binary -0.0335  -0.0396     Balanced, <0.1
## black     Binary  1.0129   0.8814 Not Balanced, >0.1
## age97    Contin.  0.2181   0.1993 Not Balanced, >0.1
## pcged97  Contin. -1.0082  -0.8629 Not Balanced, >0.1
## mratio96 Contin. -2.1090  -1.7335 Not Balanced, >0.1
##
## Balance tally for mean differences
##                    count
## Balanced, <0.1         1
## Not Balanced, >0.1     4
##
## Variable with the greatest mean difference
##  Variable Diff.Adj        M.Threshold
##  mratio96  -1.7335 Not Balanced, >0.1
##
## Effective sample sizes
##            Control Treated
## Unadjusted  729.       274
## Adjusted    663.65     274
```

3 Day 3 - September 17, 2022

3.1 Matching Estimators

3.1.1 Load Packages

A variety of matching estimators are implemented in the `Matching` package. Unlike the `MatchIt` package, which uses "subset selection" to arrive at a weighted subset of the units from the dataset, the `Matching` package uses "matching imputation" to impute potential outcomes using the observed outcomes of paired units.[7]

Because the assumptions about constant treatment effect and homoskedasticity may not be valid for certain types of data, we will use the `lmtest` package for the Breusch-Pagan Test to check this assumption.

```
library(Matching)
```

Note that the `Matching::Match()` function is intended to be used in conjunction with the `MatchBalance()` function. However, functions from the `cobalt` package also work well and tend to be cleaner in presentation:[8]

```
data(lalonde)
ex_f <- as.formula(treat ~ age + I(age^2) + educ + I(educ^2) + black +
  hisp + married + nodegr + re74 + I(re74^2) + re75 + I(re75^2) +
  u74 + u75)
ex_m1 <- glm(ex_f, family = binomial, data = lalonde)
ex_p <- ex_m1$fitted.values
ex_X <- ex_m1$fitted
ex_Y <- lalonde$re78
ex_Tr <- lalonde$treat
ex_rr <- Match(Y = ex_Y, Tr = ex_Tr, X = ex_X, M = 1)
summary(ex_rr)
ex_mb <- MatchBalance(treat ~ age + I(age^2) + educ + I(educ^2) + black +
  hisp + married + nodegr + re74 + I(re74^2) + re75 + I(re75^2) +
  u74 + u75, data = lalonde, match.out = ex_rr, nboots = 10)
(ex_bal <- cobalt::bal.tab(ex_rr, ex_f,
  data = lalonde, distance = ~ex_p, un = T,
  binary = "std", threshold = .1
))
cobalt::love.plot(ex_bal)
```

3.1.2 Description of Dataset

This example uses the 1997 Child Development Supplement (CDS) to the Panel Study of Income Dynamics (PSID) and the core PSID annual data from 1968 to 1997.

The dependent variable in this dataset is `pcss97`, a passage comprehension score. Higher scores on this measure indicate higher academic achievement. The treatment variable is `kuse` or children who ever used Aid to Families With Dependent Children (AFDC). The covariates or matching variables are:

- `pcg_adc`: Caregiver's History of Using Welfare (Number of Years; range: 0-7)
- `age97`: Child's Age in 1997
- `mratio96`: Ratio of Family Income to Poverty Line in 1996
- `pcged97`: Caregiver's Education in 1997 (Years of Schooling)
- `male`: Child's Gender: Male (1 = Male; 0 = Female)
- `black`: Child's Race: African American (1 = African American; 0 = Other)

[7] https://kosukeimai.github.io/MatchIt/articles/matching-methods.html

[8] https://cran.r-project.org/web/packages/cobalt/vignettes/cobalt_A1_other_packages.html#using-bal.tab-with-matching

3.1.3 Load Data

```
me_df <- haven::read_dta("data/cds_pcss97.dta") %>%
  haven::zap_formats() %>%
  sjlabelled::remove_all_labels() %>%
  as_tibble()
head(me_df) %>%
  kbl(booktabs = T, linesep = "", digits = 2) %>%
  kable_styling(position = "center") %>%
  kable_styling(latex_options = c("striped", "hold_position"))
```

kid	pcg_id	age97	pcss97	k_adc	pcged97	pcg_adc	mratio96	black	male	kuse
4180	4179	12	93	7.69	9	0	0.41	0	0	1
5032	5170	12	122	0.00	12	0	5.00	0	0	0
7041	7030	6	93	0.00	11	2	1.07	0	1	0
10033	10177	7	101	8.33	12	0	2.40	1	0	1
10034	10006	6	127	0.00	14	0	1.00	1	0	0
14030	14172	11	103	0.00	12	0	3.77	0	1	0

3.1.4 Breusch-Pagan Test for Heteroskedasticity

The homoskedastic variance estimator assumes that the unit-level treatment effect is constant and that the conditional variance of $Y_i(w)$ given X_i does not vary with either covariates or the treatment.

To carry out the Breusch-Pagan Test, first we regress the outcome variable on the matching variables using OLS:

```
# Regress outcome on treatment and matching variables using OLS
me_m1 <- lm(pcss97 ~ kuse + male + black + age97 + pcged97 + mratio96 + pcg_adc,
  data = me_df
)
```

Next, we can run the Breusch-Pagan test for each matching variable:

```
lmtest::bptest(me_m1, ~kuse, data = me_df, studentize = F)
lmtest::bptest(me_m1, ~male, data = me_df, studentize = F)
lmtest::bptest(me_m1, ~black, data = me_df, studentize = F)
lmtest::bptest(me_m1, ~age97, data = me_df, studentize = F) # significant
lmtest::bptest(me_m1, ~pcged97, data = me_df, studentize = F)
lmtest::bptest(me_m1, ~mratio96, data = me_df, studentize = F)
lmtest::bptest(me_m1, ~pcg_adc, data = me_df, studentize = F)
```

Or use a function to test every variable:

```
bp <- function(var, df, md) {
  lmtest::bptest(md, as.formula(paste0("~", var)), data = df, studentize = F) %>%
    broom::tidy() %>%
    mutate(variable = var) %>%
    select(variable, statistic, p.value)
}
map_dfr(c("kuse", "male", "black", "age97", "pcged97", "mratio96", "pcg_adc"), bp,
  df = me_df, md = me_m1
) %>%
  kbl(
    booktabs = T, linesep = "", digits = 2,
```

```
    caption = "Results of Breusch-Pagan Tests for Heteroskedasticity"
  ) %>%
  kable_styling(position = "center") %>%
  kable_styling(latex_options = c("striped", "hold_position"))
```

Table 1: Results of Breusch-Pagan Tests for Heteroskedasticity

variable	statistic	p.value
kuse	0.01	0.92
male	3.07	0.08
black	1.43	0.23
age97	25.27	0.00
pcged97	0.19	0.66
mratio96	0.62	0.43
pcg_adc	0.29	0.59

Results from the Breusch-Pagan tests showed that the homoskedasticity assumption is not valid for child's age (`age97`) ($p < .05$) and indicated that the conditional variance of the outcome variable was not constant across levels of child's age. Based on this finding, we should use the robust variance estimator that allows for heteroskedasticity (i.e., the `Var.calc` argument in the `Matching::Match()` function).

3.1.5 Matching Estimators

Of the six matching variables in this example, four are continuous and two are categorical, therefore bias-corrected matching estimator is necessary to correct for bias corresponding to the matching discrepancies between matched units and their matches on the four continuous covariates. Regression adjustment can be used with the `BiasAdjust = T` argument. (Tip: When you have one or more continuous covariate in your matching, always use the bias-corrected matching estimator.)

By default, the `Matching::Match()` function uses the matching variables to make bias adjustments. However, these covariates can be specified using the `Z` argument (example shown below).

The `M` argument specifies the number of matches which should be found. The default is one-to-one matching (i.e., M = 1). Abadie and Imbens suggest that M be small, and `M = 4` typically performs well in terms of mean-squared error.

If `Var.calc = 0`, then homoskedasticity is assumed. Use `Var.calc = 4` to request the robust variance estimator using four matches. This algorithm developed by Abadie and Imbens (2002) includes a second matching procedure such that it matches treated units to treated units and control units to controls.

The `estimand` argument is by default "ATT" but can be set to "ATE" or "ATC". Typically, we are interested in the ATE or ATT. The ATT seeks to answer questions such as "How would treated patients' outcomes differ, on average, had they not received treatment?"; the ATE is the treatment effect for both control and treated units; and the ATC, the average treatment effect for the controls, investigates the effect of the treatment for a population who did not receive treatment.[9]

The `sample` argument is a logical flag indicating whether the population or sample variance should be estimated. An example may help illustrate the difference between PATE and SATE: "While the SATE is useful for judging how a job-training program has affected a particular group of participants, the PATE can be used to evaluate whether another group of participants drawn from the same population is likely to benefit from the program."[10] In other words, the sample effect shows whether the program is successful in the sample at hand, while the population effect shows whether the same program would be successful in a second sample from the population.

[9] https://arxiv.org/abs/2106.10577
[10] https://journals.sagepub.com/doi/pdf/10.1177/1536867X0400400307

Results from the `Matching::Match()` function are identical to Stata's `nnmatch` program.

3.1.6 Define Outcome (Y), Treatment Index (Tr), and Variables to Match On (X)

Use the `Y`, `Tr`, and `X`, arguments in the `Match()` function to specify the outcome vector, the treatment vector, and the matrix of variables to match on, respectively. The X matrix may or may not contain the propensity score.

```
me_Y <- me_df$pcss97
me_Tr <- me_df$kuse
me_X <- select(me_df, male, black, age97, pcged97, mratio96, pcg_adc)
```

3.1.7 Get Estimators Individually

Note that by default matching is done with replacement. However, this can be changed with the `replace` argument. A matched dataset can be recovered by using the `index.treated` and `index.control` vectors. Further, `index.dropped` contains observations that have been dropped (e.g., due to a caliper setting).

```
# Sample Average Treatment Effect (SATE)
me1 <- Match(
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATE", sample = T
)
summary(me1)
head(me1$index.treated)
head(me1$index.control)
head(me1$index.dropped)
head(me1$se) # Abadie-Imbens standard error

# Population Average Treatment Effect (PATE)
summary(Match(
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATE", sample = F
))

# Sample average treatment effect for the treated (SATT)
summary(Match(
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATT", sample = T
))

# Population average treatment effect for the treated (PATT)
summary(Match(
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATT", sample = F
))

# Sample average treatment effect for the controls (SATC)
summary(Match(
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATC", sample = T
))

# Population average treatment effect for the controls (PATC)
summary(Match(
```

```
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATC", sample = F
))
```

3.1.8 Get All Estimators

```
# Function for extracting estimate, SE, t-stat, and p-value from Match()
get_match <- function(estimand, sample, Y, Tr, X) {
  m <- Matching::Match(
    Y = Y, Tr = Tr, X = X, M = 4, BiasAdjust = T, Var.calc = 4,
    estimand = estimand, sample = sample
  )
  return(list(
    est = m$est[, 1],
    se = m$se,
    t.stat = m$est[, 1] / m$se,
    p = (1 - pnorm(abs(m$est[, 1] / m$se))) * 2
  ))
}

# Estimate different matching estimators
tribble(
  ~estimator, ~estimand, ~sample,
  "SATE", "ATE", T,
  "PATE", "ATE", F,
  "SATT", "ATT", T,
  "PATT", "ATT", F,
  "SATC", "ATC", T,
  "PATC", "ATC", F
) %>%
  rowwise() %>%
  mutate(match = list(get_match(estimand, sample, me_Y, me_Tr, me_X))) %>%
  unnest_wider(match) %>%
  select(-estimand, -sample) %>%
  kbl(
    booktabs = T, linesep = "",
    caption = "Bias-Corrected Matching Estimators with Robust Standard Errors"
  ) %>%
  kable_styling(position = "center") %>%
  kable_styling(latex_options = c("striped", "hold_position"))
```

Table 2: Bias-Corrected Matching Estimators with Robust Standard Errors

estimator	est	se	t.stat	p
SATE	-4.703773	1.769696	-2.657956	0.0078616
PATE	-4.703773	1.765187	-2.664746	0.0077047
SATT	-5.229651	1.781217	-2.935999	0.0033248
PATT	-5.229651	1.720590	-3.039451	0.0023701
SATC	-4.467254	2.133536	-2.093827	0.0362754
PATC	-4.467254	2.135647	-2.091757	0.0364602

The results suggest that childhood poverty strongly affected children's academic achievement.

ATE: On average, children who used AFDC in childhood had a passage comprehension score 4.7 units lower than that of children who had never used AFDC in childhood. This effect is statistically significant in the sample at hand (SATE $p < .05$) as well as in a second sample drawn from the same population (PATE $p < .05$).

ATT: With regard to the subpopulation of treated participants, on average, children who used AFDC in childhood had a passage comprehension score 5.2 units lower than that of children who had never used AFDC in childhood. This effect is statistically significant in the sample at hand (SATT $p < .05$) as well as in a second sample drawn from the same population (PATT $p < .05$).

ATC: Had all controls (i.e., children who never used AFDC) used AFDC and all treated children had not used AFDC, then on average, the control children would have a passage comprehension score 4.5 units lower than their counterparts (SATC $p < .05$; PATC $p < .05$).

Additional observations:

1. A population effect indicates whether the tested intervention will be effective in a second sample taken from the same population. Taking SATT and PATT as examples, the study indicated that the treatment effect for the treated group was statistically significant in the sample at a level of .01. If we take a second sample from the population, we are likely to observe the same level of treatment effect for the treated, and the effect should remain statistically significant at a level of .01. The point estimate of the population effect is identical to the point estimate of its corresponding sample effect. A population effect differs from its corresponding sample effect on variance, and thus significance test on a population effect may have a different conclusion than that on its corresponding sample effect. Note that if treated units are discarded (e.g., due to a caliper), the estimand no longer estimates a particular population but rather only the average treatment effect on the remaining matched units (ATM).
2. Note that in this study, SATT = -5.23 and SATC = -4.47, or a difference of 0.76 units. This difference is attributable either to additional selection bias that was not accounted for in the study or to study data that violated assumptions of matching estimators, which suggests the need for further scrutiny.
3. All six treatment effects were statistically significant ($p < .05$). Thus, we can conclude that the study data could not reject a null hypothesis of a zero treatment effect, and childhood poverty appears to be an important factor causing children's poor achievement in passage comprehension.

3.1.9 Specify Variables in the Bias-Corrected Matching

The `Z` argument can be used to specify the covariates for which we wish to make bias adjustments.

```
# Sample Average Treatment Effect (SATE)
me_Z <- select(me_df, age97, pcged97, mratio96, pcg_adc)
summary(Matching::Match(
  Y = me_Y, Tr = me_Tr, X = me_X, Z = me_Z, M = 4,
  BiasAdjust = T, Var.calc = 4, estimand = "ATE", sample = T
))
```

```
##
## Estimate...  -4.4867
## AI SE......  1.7697
## T-stat.....  -2.5353
## p.val......  0.011235
##
## Original number of observations..............  606
## Original number of treated obs...............  188
## Matched number of observations...............  606
## Matched number of observations  (unweighted).  2441
```

3.1.10 Check Balance

The `Matching::Match()` function works well in conjunction with the `cobalt::bal.tab()` function for checking covariate balance.

By default, the denominator for standardized mean differences uses a pooled estimate (square root of the average of the group variances) for ATE and the standard deviation of the treated group for ATT, and both standard deviations are computed using the sample before matching. This option can also be manually set with the `s.d.denom` option.

```
# Example of Checking Balance for SATE
me_SATE <- Match(
  Y = me_Y, Tr = me_Tr, X = me_X, M = 4, BiasAdjust = T, Var.calc = 4,
  estimand = "ATE", sample = T
)
(me_SATE_bal <- cobalt::bal.tab(
  me_SATE, kuse ~ male + black + age97 + pcged97 + mratio96 + pcg_adc,
  data = me_df,
  abs = T,
  un = T,
  binary = "std",
  thresholds = c(m = .1),
  s.d.denom = "pooled"
))
```

```
## Balance Measures
##             Type Diff.Un Diff.Adj        M.Threshold
## male      Binary  0.0480   0.0099     Balanced, <0.1
## black     Binary  0.9738   0.0553     Balanced, <0.1
## age97    Contin.  0.0713   0.0010     Balanced, <0.1
## pcged97  Contin.  0.8084   0.3478 Not Balanced, >0.1
## mratio96 Contin.  1.0066   0.5831 Not Balanced, >0.1
## pcg_adc  Contin.  0.6508   0.0450     Balanced, <0.1
##
## Balance tally for mean differences
##                    count
## Balanced, <0.1         4
## Not Balanced, >0.1     2
##
## Variable with the greatest mean difference
##  Variable Diff.Adj        M.Threshold
##  mratio96   0.5831 Not Balanced, >0.1
##
## Sample sizes
##                      Control Treated
## All                   418.    188.
## Matched (ESS)         318.46   85.38
## Matched (Unweighted)  418.    188.
```

```
cobalt::love.plot(me_SATE_bal) +
  labs(title = "Covariate Balance for SATE")
```

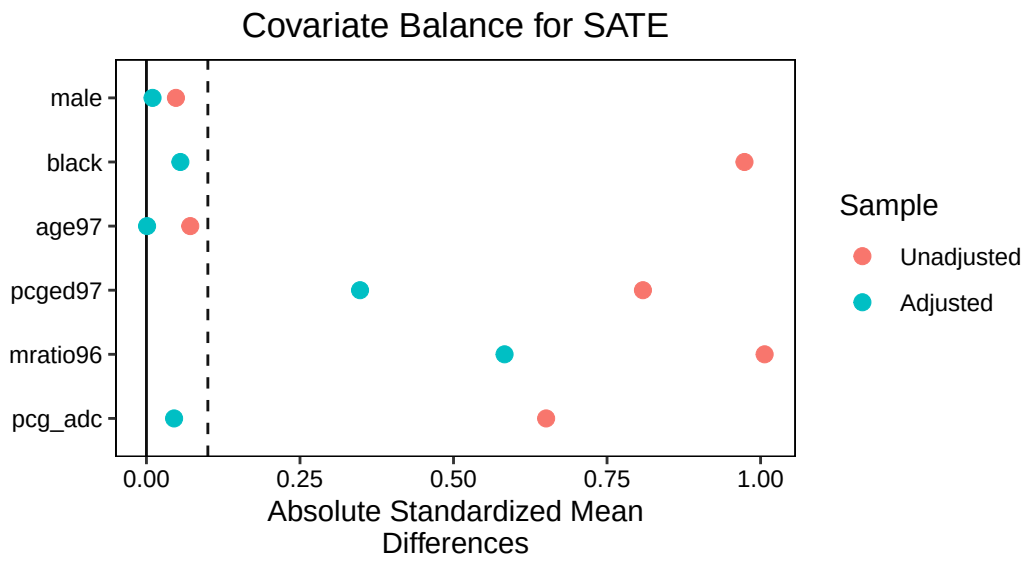
Covariate Balance for SATE
male
black
age97
pcged97
mratio96
pcg_adc
0.00
0.25
0.50
0.75
1.00
Absolute Standardized Mean Differences
Sample
Unadjusted
Adjusted

4 Appendix A: IMBALANCE Stata Module in R

```
imbalance <- function(df, varname, treatname, blockname) {

  # dx
  df2 <- df %>%
    group_by({{ treatname }}) %>%
    summarise(
      m_x = mean({{ varname }}),
      sd_x = sd({{ varname }}),
      .groups = "drop"
    )
  mxt <- df2[2, 2]
  mxc <- df2[1, 2]
  s2xt <- df2[2, 3]^2
  s2xc <- df2[1, 3]^2
  sx <- sqrt((s2xt + s2xc) / 2)
  dx <- as.numeric(abs(mxt - mxc) / sx)

  # dx
  df3 <- df %>%
    group_by({{ blockname }}, {{ treatname }}) %>%
    summarise(
      m_x = mean({{ varname }}),
      sd_x = sd({{ varname }}),
      n = n(),
      .groups = "drop"
    )

  mxc <- as.numeric(mean(filter(df3, {{ treatname }} == 0)$m_x))
  mxt <- as.numeric(mean(filter(df3, {{ treatname }} == 1)$m_x))
  dxm_num <- abs(mxt - mxc)
  dxm <- as.numeric(dxm_num / sx)

  return(list(dx = dx, dxm = dxm))
}
```

4.1 Example Usage

```
imbalance(df, mratio96, kuse, fm)
```

5 Appendix B: HODGESL Stata Module in R

```
hodgesl <- function(dataname, varname, blockname, treatname) {
  blockname_str <- deparse(substitute(blockname))
  set.seed(1000)
  renamed_file <- dataname %>%
    filter(!is.na({{ blockname }}))
  r1 <- renamed_file %>%
    group_by({{ blockname }}) %>%
    summarise(m_y = mean({{ varname }}), .groups = "drop") %>%
    arrange({{ blockname }})
  r2 <- renamed_file %>%
    group_by({{ blockname }}, {{ treatname }}) %>%
    summarise(
      mean_y = mean({{ varname }}),
      n = n(), .groups = "drop"
    ) %>%
    mutate(mean_diff = ifelse({{ treatname }} == 1,
      ((n + lag(n)) / sum(n)) * (mean_y - lag(mean_y)), NA
    )) %>%
    mutate(tx_effect = sum(mean_diff, na.rm = T)) %>%
    mutate(i = row_number()) %>%
    slice(1) %>%
    select(tx_effect, i)
  fm_results <- renamed_file %>%
    group_by({{ blockname }}, {{ treatname }}) %>%
    summarise(
      mean_y = mean({{ varname }}),
      n = n(), .groups = "drop"
    )
  r3 <- renamed_file %>%
    group_by({{ blockname }}, {{ treatname }}) %>%
    summarise(m_or_n = n(), .groups = "drop") %>%
    arrange({{ blockname }}, {{ treatname }}) %>%
    mutate(mi = ifelse({{ treatname }} == 0, m_or_n,
      ifelse({{ treatname }} == 1, NA, NA)
    )) %>%
    mutate(ni = ifelse({{ treatname }} == 1, m_or_n,
      ifelse({{ treatname }} == 0, NA, NA)
    )) %>%
    mutate(Ni = ni + lag(mi)) %>%
    mutate(mi = ifelse(is.na(mi), lag(mi), mi)) %>%
    filter(!is.na(Ni)) %>%
    mutate(factor = (mi * ni) / (Ni * (Ni - 1))) %>%
    select({{ blockname }}, factor) %>%
    arrange({{ blockname }})
  r4 <- renamed_file %>%
    arrange({{ blockname }}) %>%
    left_join(r1, by = blockname_str) %>%
    mutate(dy = {{ varname }} - m_y) %>%
    arrange(dy) %>%
    mutate(rk = row_number()) %>%
    arrange({{ blockname }})
```

```
  r4a <- r4 %>%
    filter({{ treatname }} != 0) %>%
    group_by({{ blockname }}) %>%
    summarise(wsi = sum(rk), .groups = "drop")
  r5 <- r4 %>%
    group_by({{ blockname }}) %>%
    summarise(ki_ = mean(rk), .groups = "drop")
  r6 <- r4 %>%
    filter({{ treatname }} != 0) %>%
    group_by({{ blockname }}) %>%
    summarise(ni = n(), .groups = "drop") %>%
    arrange({{ blockname }}) %>%
    left_join(r5, by = blockname_str) %>%
    mutate(E_wsi = ni * ki_) %>%
    arrange({{ blockname }})
  r7 <- r4 %>%
    arrange({{ blockname }}) %>%
    left_join(r5, by = blockname_str) %>%
    mutate(k = (rk - ki_)^2) %>%
    group_by({{ blockname }}) %>%
    summarise(ss_kd_i = sum(k), .groups = "drop") %>%
    arrange({{ blockname }})
  results <- r3 %>%
    arrange({{ blockname }}) %>%
    left_join(r7, by = blockname_str) %>%
    left_join(r6, by = blockname_str) %>%
    left_join(r4a, by = blockname_str) %>%
    mutate(
      var_wsi = factor * ss_kd_i,
      var = sum(var_wsi),
      sum_Ewsi = sum(E_wsi),
      ws = sum(wsi),
      HL_mean = ws - sum_Ewsi,
      HL_se = sqrt(var),
      z = HL_mean / HL_se,
      p = 1 - pnorm(abs(z))
    ) %>%
    select(HL_mean, HL_se, z, p) %>%
    slice(1) %>%
    mutate(i = row_number()) %>%
    left_join(r2, by = "i")
  return(results)
}
```